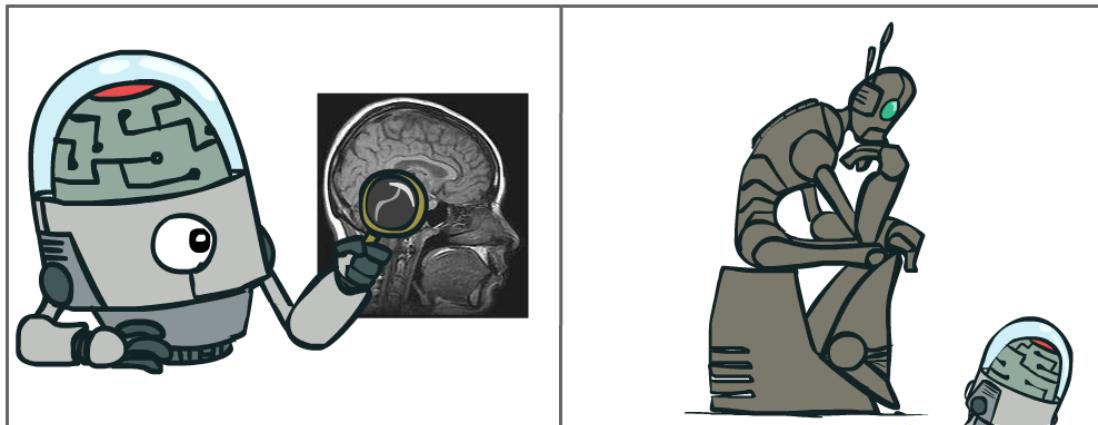


What is AI?

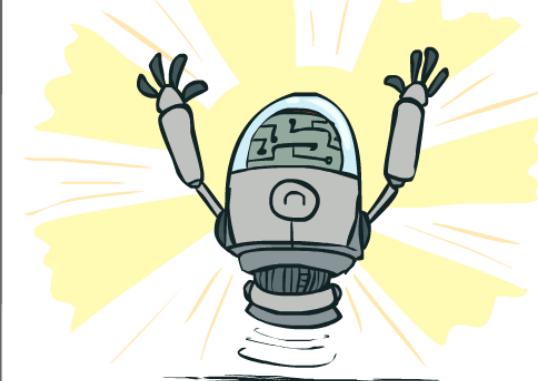
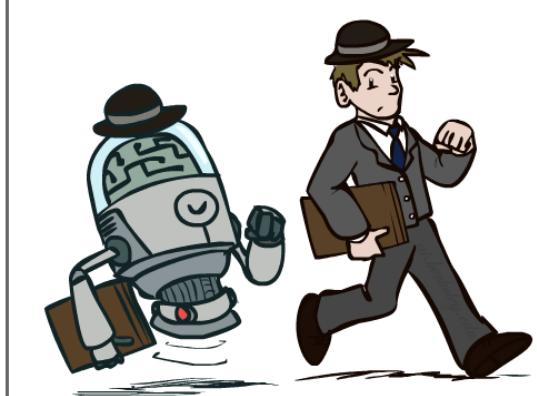
The science of making machines that:

Think like people



Think rationally

Act like people



Act rationally

Fundamental question for this lecture
(and really this whole AI field!):

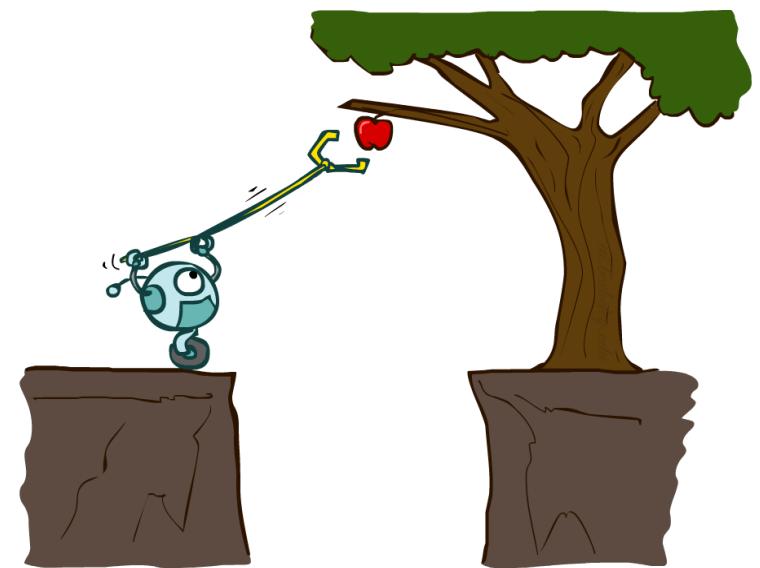
**How do you turn a real-world
problem into an AI solution?**

AI – Agents and Environments

Much (though not all!) of AI is concerned with **agents** operating in **environments**.

Agent – an entity that *perceives* and *acts*

Environment – the problem setting



Fleshing it out

Performance – measuring desired outcomes

Environment – what populates the task's world?

Actuators – what can the agent act with?

Sensors – how can the agent perceive the world?



PEAS in a taxi

Automated taxi driver

Performance – Safe, fast, legal, comfortable trip, maximize profits

Environment – Roads, other traffic, pedestrians, customers

Actuators – Steering, accelerator, brake, signals, horn, display

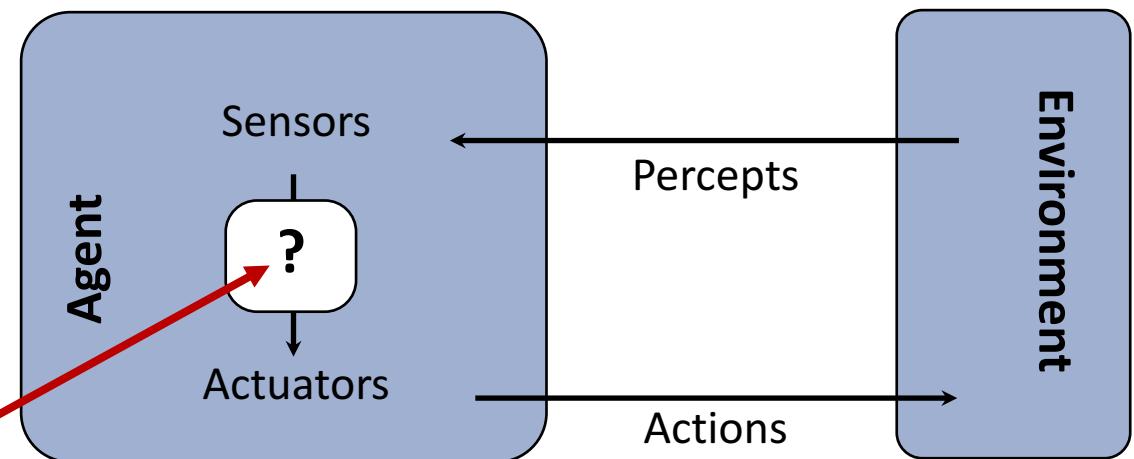
Sensors - Cameras, sonar, speedometer, GPS, odometer,
accelerometer, engine sensors, microphone/keyboard

What makes an Agent?

Agent – an entity that perceives its environment through sensors, and acts on it with actuators.

Percepts are constrained by
Sensors + Environment

Actions are constrained by
Actuators + Environment



Agent Function – how does it
choose the action?

What makes one rational?

Actually pretty simple:

**A rational agent always acts to
maximize its expected performance
measure, given current state/percept**

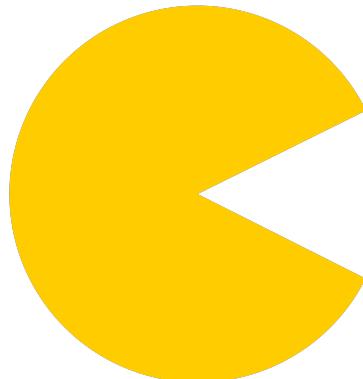
Our sample agents

Pacman

Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts



Spam detector

Percepts – sender, subject line, body of current email

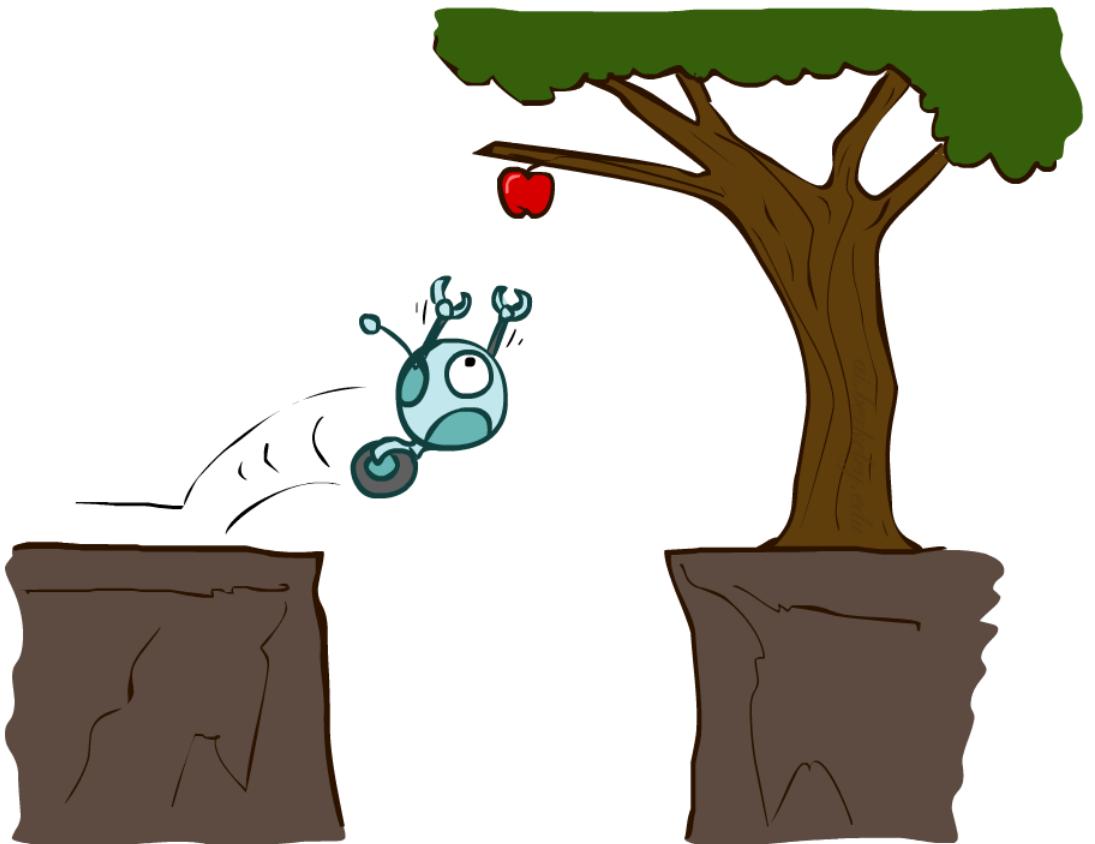
Actions – mark Spam/Not Spam

Environment – your email inbox



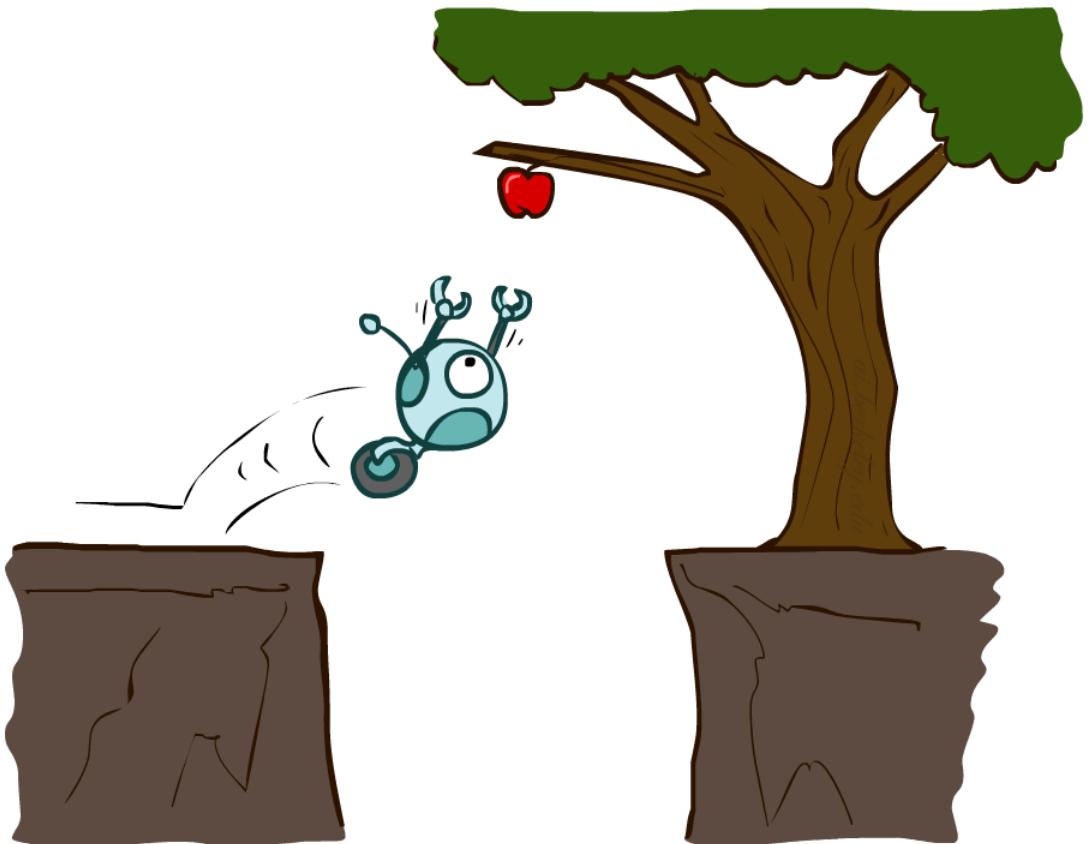
Reflex Agents

- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state
 - Do not consider the future consequences of their actions
 - Consider how the world IS



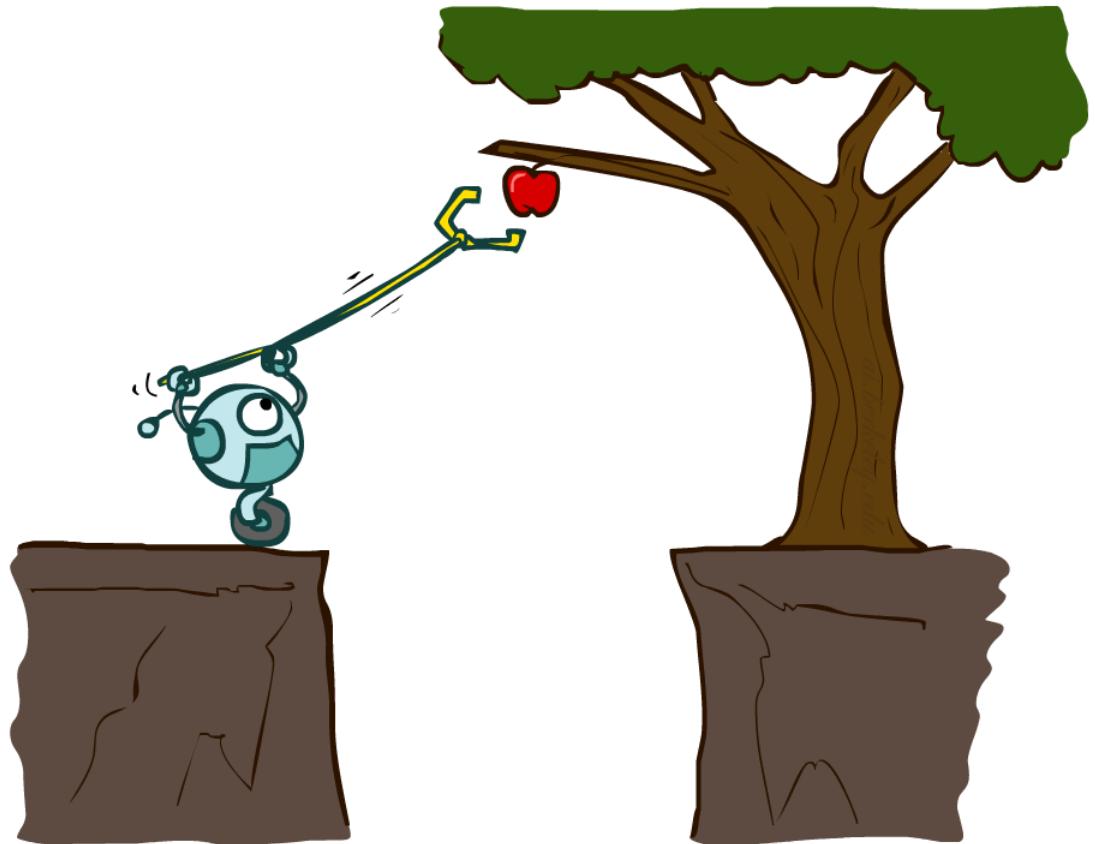
Reflex Agents

- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state
 - Do not consider the future consequences of their actions
 - Consider how the world IS
- Can a reflex agent be rational?



Planning Agents

- Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
 - Consider how the world **WOULD BE**



Goal-based Agents

Chooses action (sequence) to get from current state to some goal

Pacman

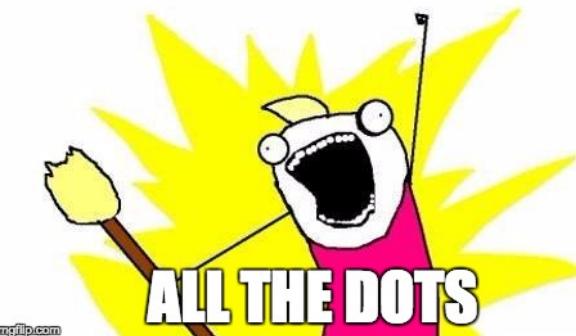
Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts

Goal:

EAT



Spam detector

Percepts – sender, subject line, body of current email

Actions – mark Spam/Not Spam

Environment – your email inbox

Goal:

???

Utility-based Agents

Chooses action (sequence) to get from current state to some goal
with maximum utility along the way

Pacman

Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts

Goal:



Spam detector

Percepts – sender, subject line, body of current email

Actions – mark Spam/Not Spam

Environment – your email inbox

Goal:



...in as short a path as possible!

Summary

Reflex agents

Act on current state (and maybe past)

Simple – current p

Model – current p

of rest of



Goal-based agents

From current state to desired future

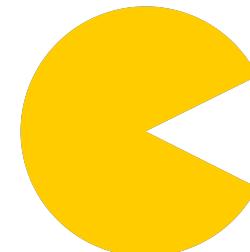
Goal only – find *any* action(s) to

reach the goal

first action(s) to

reach the goal

Can also have a **Learning Agent** –
we'll talk about these later in the
course!

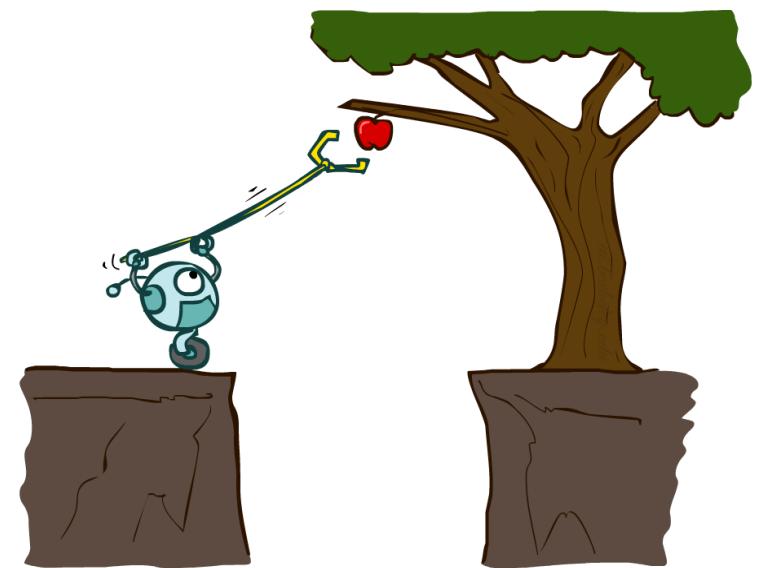


AI – Agents and Environments

Much (though not all!) of AI is concerned with **agents** operating in environments.

Agent – an entity that *perceives* and *acts*

Environment – the problem setting



Kinds of task environments

6 common properties to distinguish tasks (not exhaustive)

- **Fully observable vs Partially observable**
- **Single agent vs Multiagent**
- **Deterministic vs Stochastic**
- **Episodic vs Sequential**
- **Static vs Dynamic**
- **Discrete vs Continuous**

Fully observable vs partially observable

Fully observable – agent is able to sense everything in the environment

ACROSS

1 See 24-Across

6 They radiate outward from an earthquake's epicenter

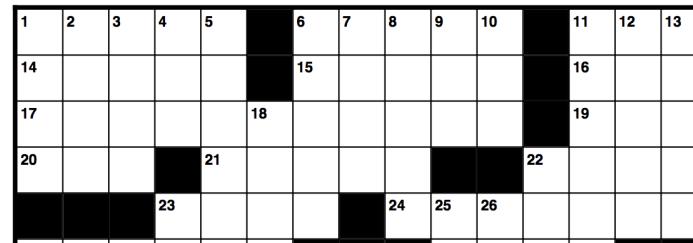
11 The "F" of "T.G.I.F.": Abbr.

45 ___ fire under (urged to take action): 2 wds.

47 Daniel Defoe's "Robinson ___"

49 Vibrations caused by earthquakes

52 Low in fat



Partially observable – noisy, inaccurate, or incomplete sensors



Single agent vs Multiagent

Single agent – self-explanatory



Multiagent – task involves more than one agent, each with its own performance measure

May be **competitive** (measures are opposed)
or **cooperative** (measures align)

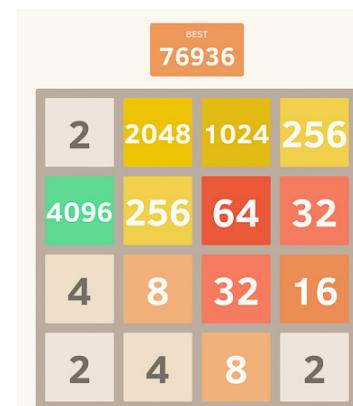


Deterministic vs Stochastic

Deterministic – next state of the world is fully determined by current state + agent action



Stochastic – it's not deterministic

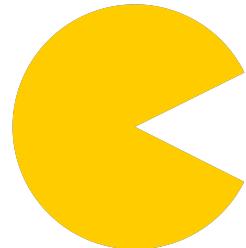


Episodic vs Sequential

Episodic – Each step/decision is independent of the previous ones



Sequential – Each step/decision affects later ones



Static vs Dynamic

Static – world doesn't change while agent is choosing an action



Dynamic – decision time matters!



Discrete vs Continuous

Discrete – possible states/actions are distinct; world changes discretely



Continuous – states/actions take on continuous values



These help determine how to approach problems

Static -> can focus on getting really high accuracy/utility

Dynamic -> trade some utility for higher efficiency (speed!)

Episodic -> reflex agent with a great model

Sequential -> need a goal-oriented agent

Stochastic -> need robustness to uncertainty/failure (robots!)

Deterministic -> can focus on efficiency and exactness (Internet crawler)

Next up

Defining search problems – how to choose the right action sequence?

Uninformed search approaches – simple reflex agents for searching