


# What is AI?


The science of making machines that:

Think like people




Think rationally

Act like people



Act rationally



Fundamental question for this lecture  
(and really this whole AI field!):

**How do you turn a real-world  
problem into an AI solution?**

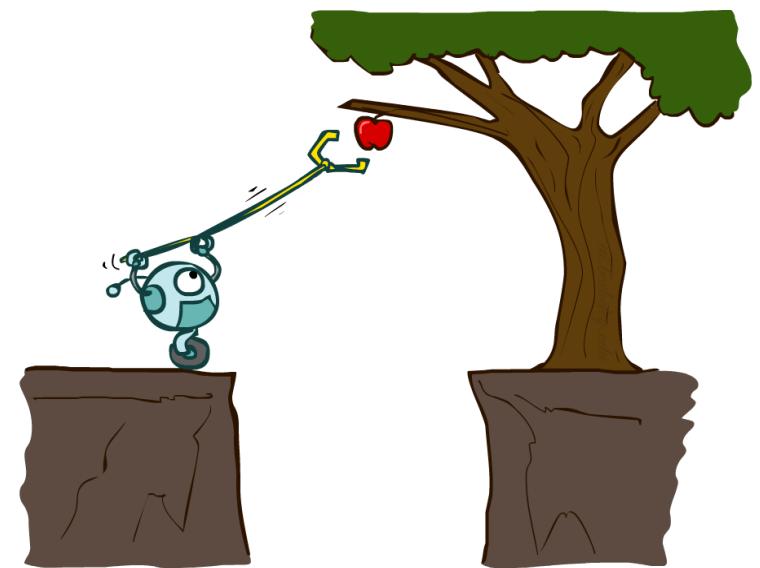
# AI – Agents and Environments

---

Much (though not all!) of AI is concerned with **agents** operating in **environments**.

**Agent** – an entity that *perceives* and *acts*

**Environment** – the problem setting



# Fleshing it out

---

**P**erformance – measuring desired outcomes

**E**nvironment – what populates the task's world?

**A**ctuators – what can the agent act with?

**S**ensors – how can the agent perceive the world?



# PEAS in a taxi

---

## Automated taxi driver

**Performance** – Safe, fast, legal, comfortable trip, maximize profits

**Environment** – Roads, other traffic, pedestrians, customers

**Actuators** – Steering, accelerator, brake, signals, horn, display


**Sensors** - Cameras, sonar, speedometer, GPS, odometer,  
accelerometer, engine sensors, microphone/keyboard

# What makes an Agent?

**Agent** – an entity that perceives its environment through sensors, and acts on it with actuators.

Percepts are constrained by  
Sensors + Environment

Actions are constrained by  
Actuators + Environment



**Agent Function** – how does it  
choose the action?

# What makes one rational?

---

Actually pretty simple:

**A rational agent always acts to  
maximize its expected performance  
measure, given current state/percept**

# Our sample agents

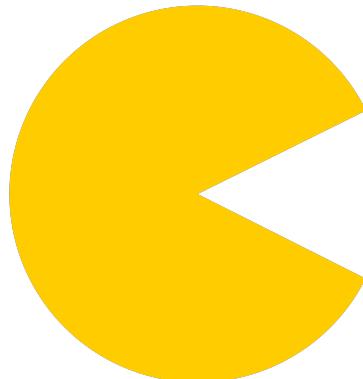
---

## Pacman

Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts



## Spam detector

Percepts – sender, subject line, body of current email

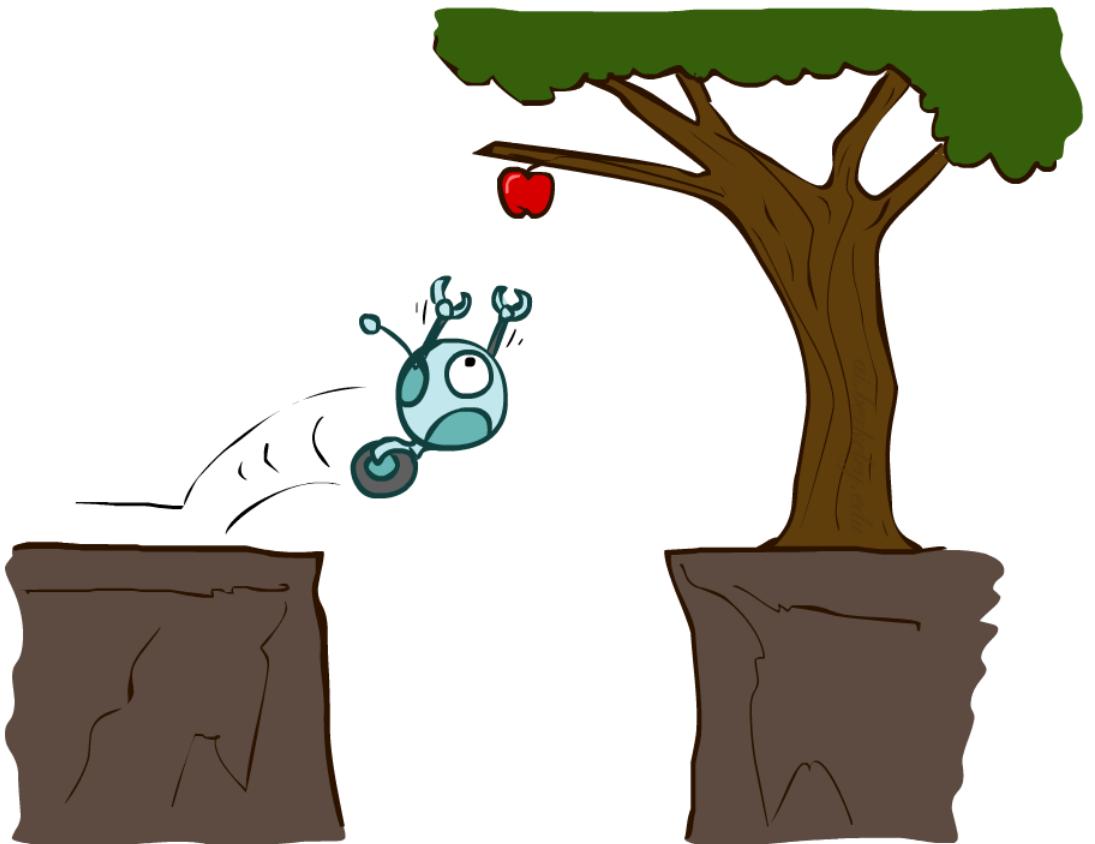
Actions – mark Spam/Not Spam

Environment – your email inbox



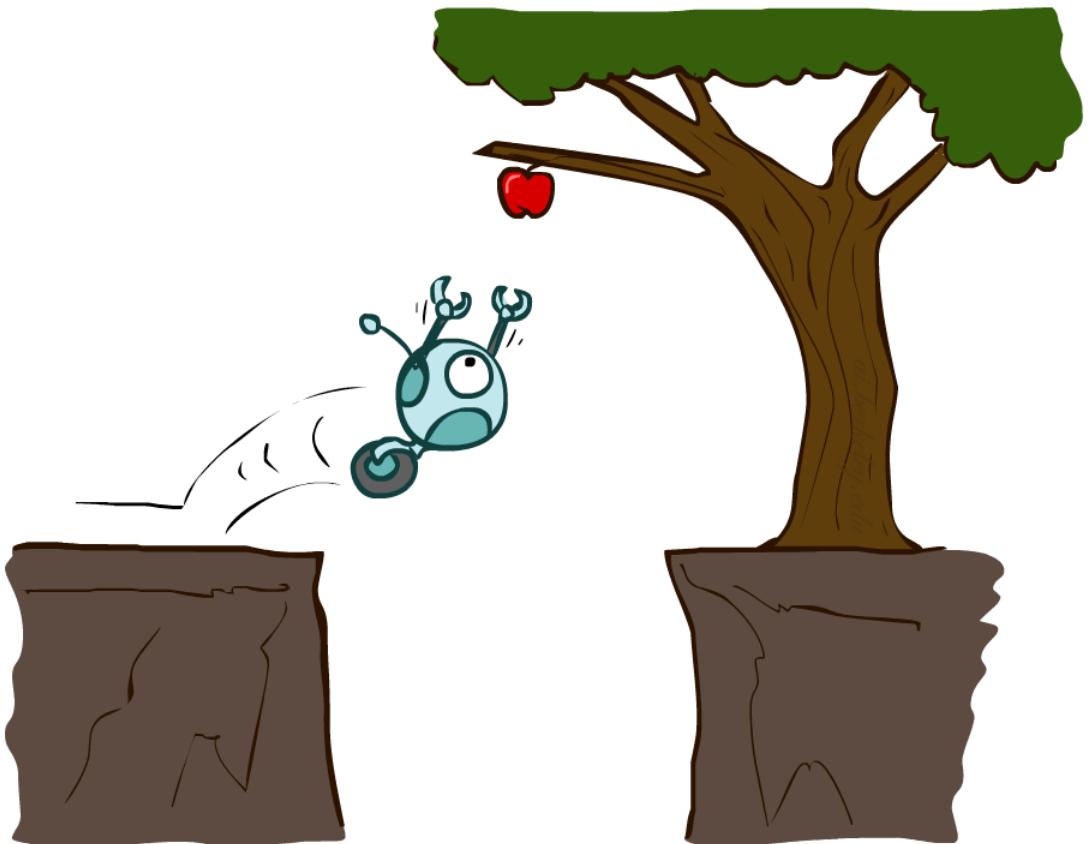
# Reflex Agents

- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS



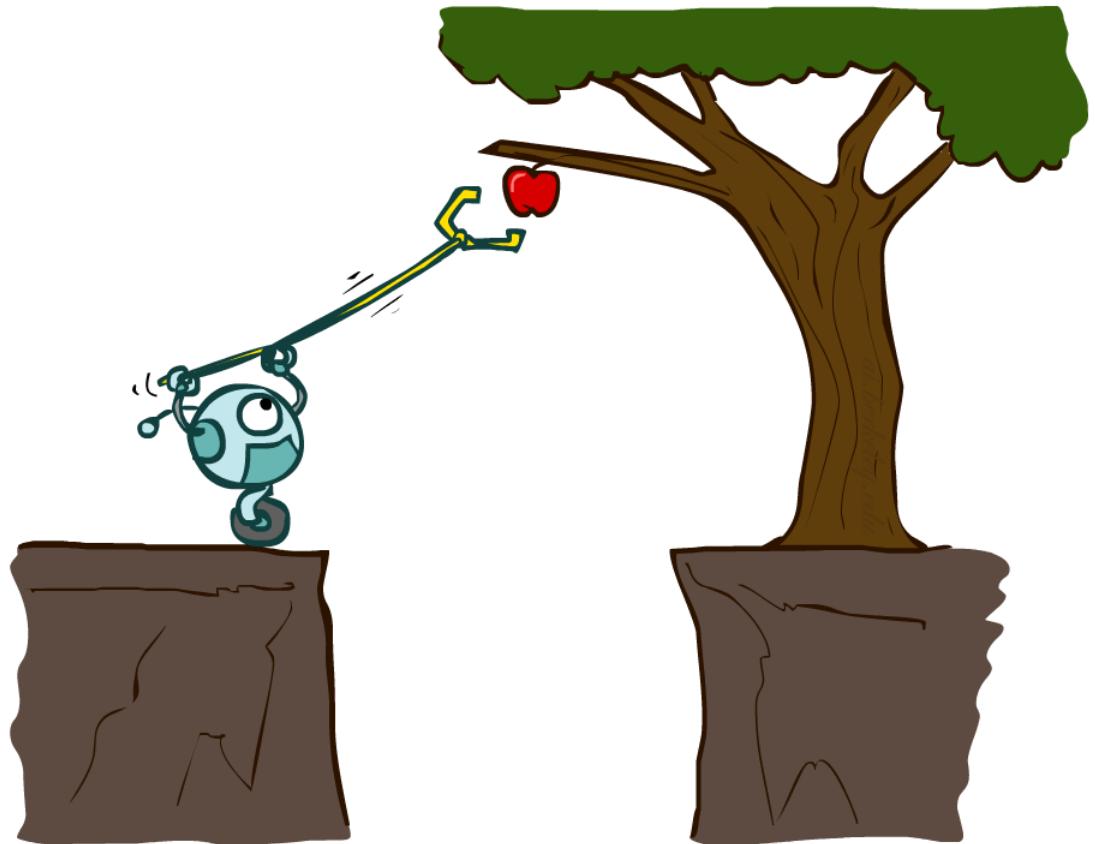
# Reflex Agents

- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- Can a reflex agent be rational?



# Planning Agents

- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - Consider how the world **WOULD BE**



# Goal-based Agents

Chooses action (sequence) to get from current state to some goal

## Pacman

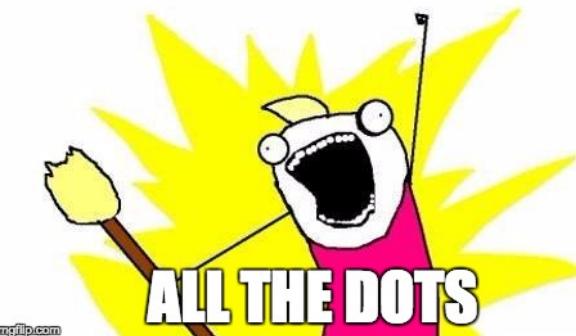
Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts

Goal:

EAT



## Spam detector

Percepts – sender, subject line, body of current email

Actions – mark Spam/Not Spam

Environment – your email inbox

Goal:

???

# Utility-based Agents

Chooses action (sequence) to get from current state to some goal  
with maximum utility along the way

## Pacman

Percepts – squares around Pacman

Actions – move U/D/L/R

Environment – map with walls, dots, and ghosts

Goal:



## Spam detector

Percepts – sender, subject line, body of current email

Actions – mark Spam/Not Spam

Environment – your email inbox

Goal:



*...in as short a path as possible!*

# Summary

## Reflex agents

Act on current state (and maybe past)

Simple – current p

Model – current p

of rest of



## Goal-based agents

From current state to desired future

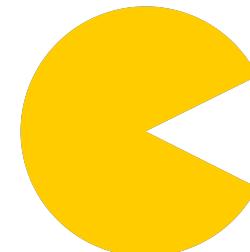
Goal only – find *any* action(s) to

reach the goal

first action(s) to

reach the goal

Can also have a **Learning Agent** –  
we'll talk about these later in the  
course!

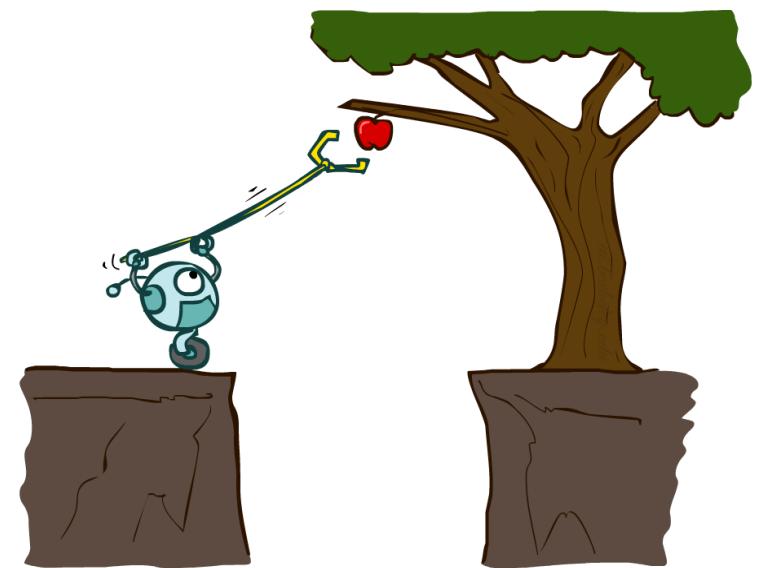


# AI – Agents and Environments

Much (though not all!) of AI is concerned with **agents** operating in environments.

**Agent** – an entity that *perceives* and *acts*

**Environment** – the problem setting



# Kinds of task environments

---

6 common properties to distinguish tasks (not exhaustive)

- **Fully observable vs Partially observable**
- **Single agent vs Multiagent**
- **Deterministic vs Stochastic**
- **Episodic vs Sequential**
- **Static vs Dynamic**
- **Discrete vs Continuous**

# Fully observable vs partially observable

**Fully observable** – agent is able to sense everything in the environment

**ACROSS**

**1** See 24-Across

**6** They radiate outward from an earthquake's epicenter

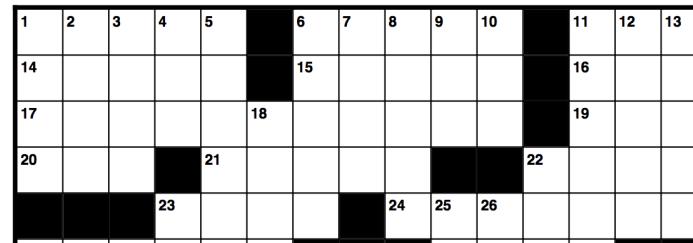
**11** The "F" of "T.G.I.F.": Abbr.

**45** \_\_\_ fire under (urged to take action): 2 wds.

**47** Daniel Defoe's "Robinson \_\_\_"

**49** Vibrations caused by earthquakes

**52** Low in fat



**Partially observable** – noisy, inaccurate, or incomplete sensors



# Single agent vs Multiagent

**Single agent** – self-explanatory



**Multiagent** – task involves more than one agent, each with its own performance measure

May be **competitive** (measures are opposed)  
or **cooperative** (measures align)




# Deterministic vs Stochastic

**Deterministic** – next state of the world is fully determined by current state + agent action



**Stochastic** – it's not deterministic



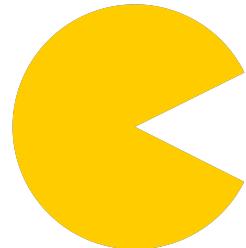
# Episodic vs Sequential

---

**Episodic** – Each step/decision is independent of the previous ones



**Sequential** – Each step/decision affects later ones



# Static vs Dynamic

---

**Static** – world doesn't change while agent is choosing an action



**Dynamic** – decision time matters!



# Discrete vs Continuous

**Discrete** – possible states/actions are distinct; world changes discretely



**Continuous** – states/actions take on continuous values



# These help determine how to approach problems

---

**Static** -> can focus on getting really high accuracy/utility

**Dynamic** -> trade some utility for higher efficiency (speed!)

**Episodic** -> reflex agent with a great model

**Sequential** -> need a goal-oriented agent

**Stochastic** -> need robustness to uncertainty/failure (robots!)

**Deterministic** -> can focus on efficiency and exactness (Internet crawler)

Next up

---

**Defining search problems** – how to choose the right action sequence?

**Uninformed search approaches** – simple reflex agents for searching