

# Announcements

---

**Homework 1** will be posted this afternoon. After today, you should be able to do Questions 1-3.

**Python 2.7 only** for the homework; turns out Python 3 plays merry hell with the autograder.

# Kinds of task environments

---

6 common properties to distinguish tasks (not exhaustive)

- **Fully observable vs Partially observable**
- **Single agent vs Multiagent**
- **Deterministic vs Stochastic**
- **Episodic vs Sequential**
- **Static vs Dynamic**
- **Discrete vs Continuous**

Col 1 – Poker

Col 2 – Self-driving taxi

Col 3 – Spam classifier

Col 4 – Pacman with ghosts

Col 5 – Oil refinery control system

Col 6 – Automatic speech transcription

# AI in practice



---

Mostly trying to define a problem in such a way  
that someone else has already solved it!

- Optimally (or close)
- Efficiently

# Planning Agents

- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - Consider how the world **WOULD BE**



# Search Problems

---



Search?

---

**Nope!**




Google Search

I'm Feeling Lucky

# Information Retrieval vs Search


Google

*(Information Retrieval)*



Search

*(Problem-Solving)*



# Definition of Search

---

**Finding a (best) sequence of actions  
to solve a problem**

For now, assume the problem is


- Deterministic
- Fully observable
- Known




# Search Problem Mechanics

- A **search problem** consists of:

- A state space



- A successor function  
(with actions, costs)




- A start state and a goal test
- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

# Search Problems Are Models

---




# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment




A **search state** keeps only the details needed for planning (abstraction)


- Problem: Pathing
  - States:  $(x,y)$  location
  - Actions: NSEW
  - Successor: update location only
  - Goal test: is  $(x,y)=END$
- Problem: Eat-All-Dots
  - States:  $\{(x,y), \text{dot booleans}\}$
  - Actions: NSEW
  - Successor: update location and possibly a dot boolean
  - Goal test: dots all false

# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for pathing?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$



# Quiz: Safe Passage




- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)

# Search Problem Mechanics


- A **search problem** consists of:

- A state space



- A successor function  
(with actions, costs)

- A start state and a goal test




- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

What are some problems that can't be formulated as search?

# State Space Graphs and Search Trees


---






# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



*Tiny search graph for a tiny search problem*


# Search Trees



- A search tree:
  - A “what if” tree of plans and their outcomes
  - The start state is the root node
  - Children correspond to successors
  - Nodes show states, but correspond to ACTION SEQUENCES that achieve those states
  - For most problems, we can never actually build the whole tree

# State Space Graphs vs. Search Trees


## State Space Graph



*Each NODE in the search tree corresponds to an entire PATH in the state space graph.*

*We construct both on demand – and we construct as little as possible.*


## Search Tree



# Quiz: State Space Graphs vs. Search Trees

---

Consider this 4-state graph:




How big is its search tree (from S)?




Important: Lots of repeated structure in the search tree!

# Tree Search


---



# Search Example: Romania



# Searching with a Search Tree



- Search:
  - Expand out potential plans (tree nodes)
  - Maintain a **fringe** of partial plans under consideration
  - Try to expand as few tree nodes as possible




# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?


# Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^{m+1})$



# Breadth-First Search


---




# Breadth-First Search

*Strategy: expand a shallowest node first*

*Implementation: Fringe is a FIFO queue*




Search  
Tiers




# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if costs are all 1 (more on costs later)




# Depth-First Search

---




*Implementation:  
Fringe is a LIFO stack*



# Depth-First Search (DFS) Properties


- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost






# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. ....
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!



# Quiz: DFS vs BFS

---




# Quiz: DFS vs BFS

---

- When will BFS outperform DFS?
- When will DFS outperform BFS?
- Pathological examples w/ 10 nodes?

# Cost-Sensitive Search





BFS finds the shortest path in terms of number of actions.  
It does not find the least-cost path. We will now cover  
a similar algorithm which does find the least-cost path.

# Uniform Cost Search


*Strategy: expand a cheapest node first:*

*Fringe is a priority queue  
(priority: cumulative cost)*





# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the “effective depth” is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (skipping the proof for now)



# Uniform Cost Issues

- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We’ll fix that soon!



Next time

---

Keep thinking: What are some problems that can't be formulated as search?

**Homework 1 going up later today**

**Informed search methods**