

# First of all what is Deep Learning?

- Composition of non-linear transformation of the data.
- Goal: **Learn useful representations, aka features, directly from data.**
- Many varieties, can be unsupervised or supervised.
- Today is about ConvNets, which is a **supervised** deep learning method.

# Recap: Supervised Learning

$\{(x^i, y^i), i=1 \dots P\}$  training dataset

$x^i$  i-th input training example

$y^i$  i-th target label

$P$  number of training examples



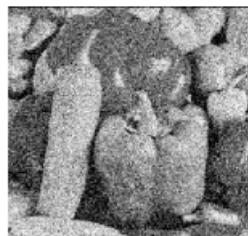
# Supervised Learning: Examples

## Classification



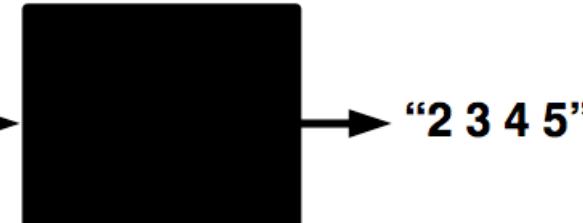
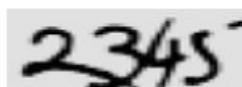
classification

## Denoising



regression

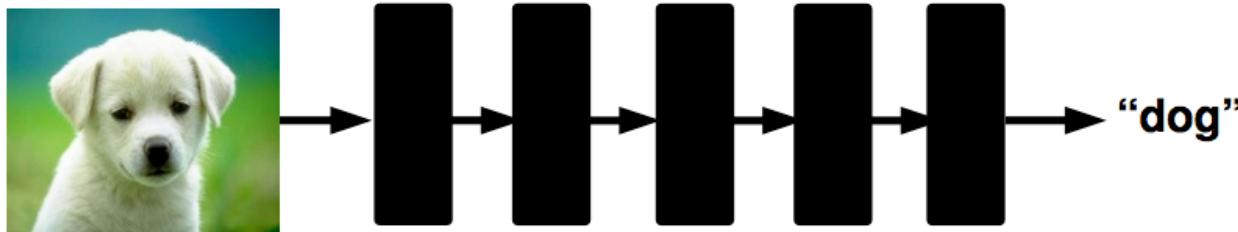
## OCR



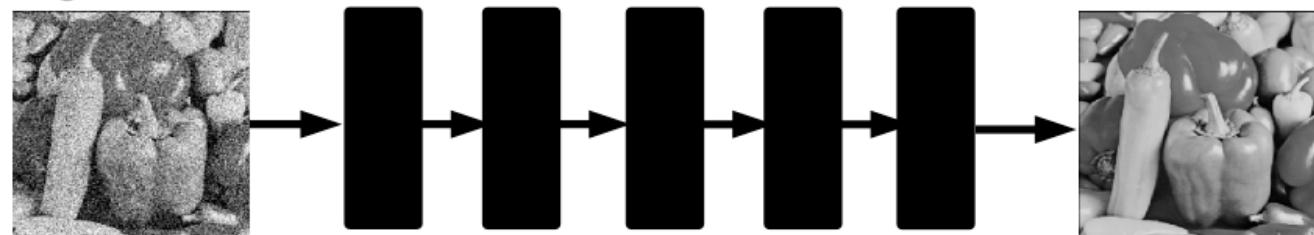
structured  
prediction

# Supervised Deep Learning

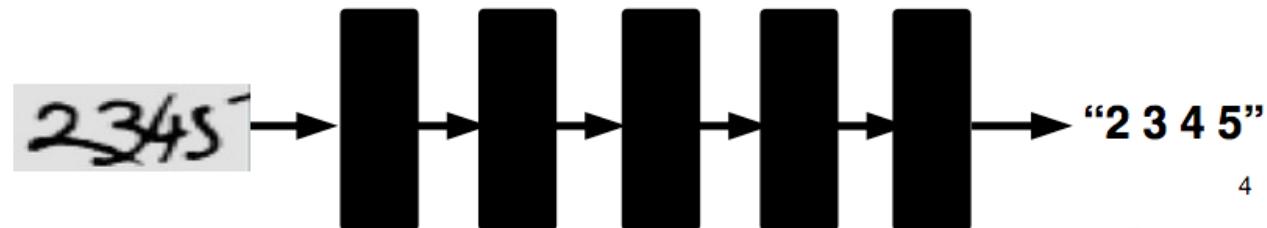
## Classification



## Denoising



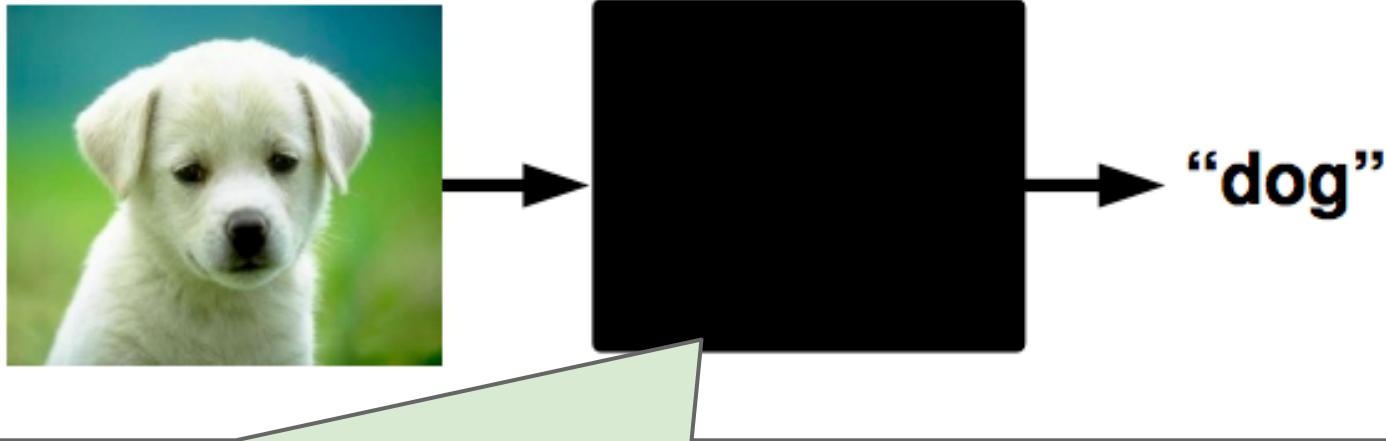
## OCR



So deep learning is about learning  
feature representation in a  
compositional manner.

But wait,  
**why learn features?**

# The Black Box in a Traditional Recognition Approach



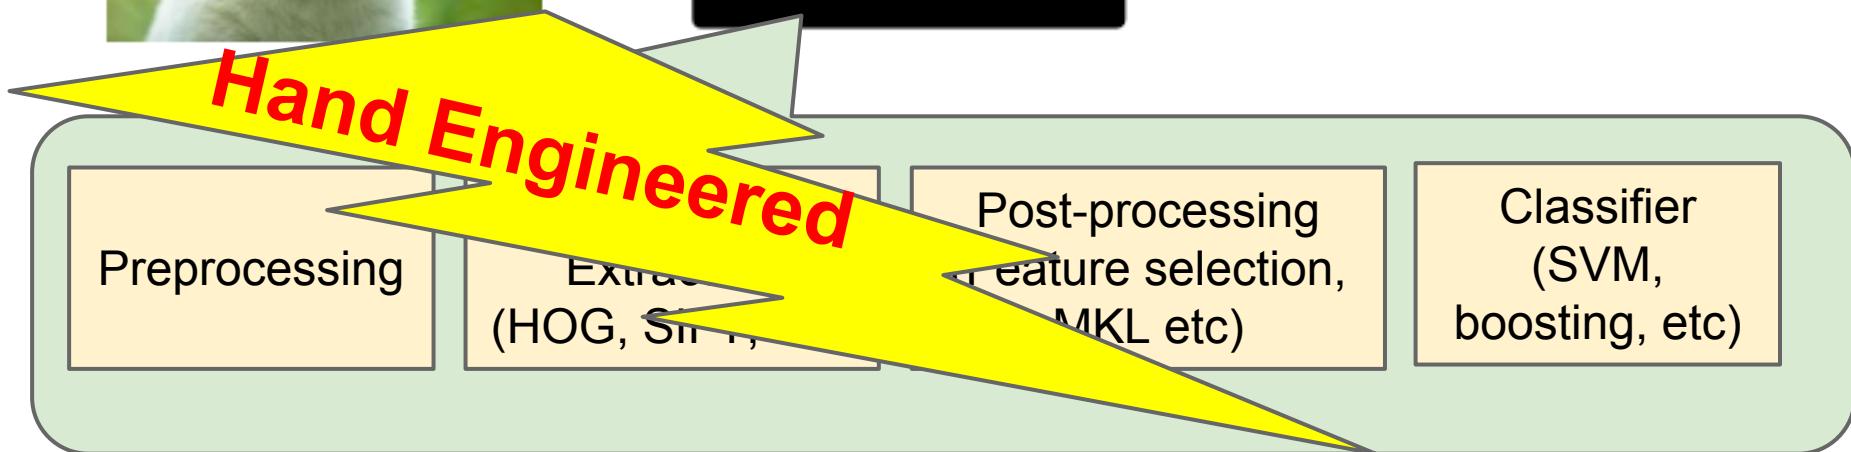
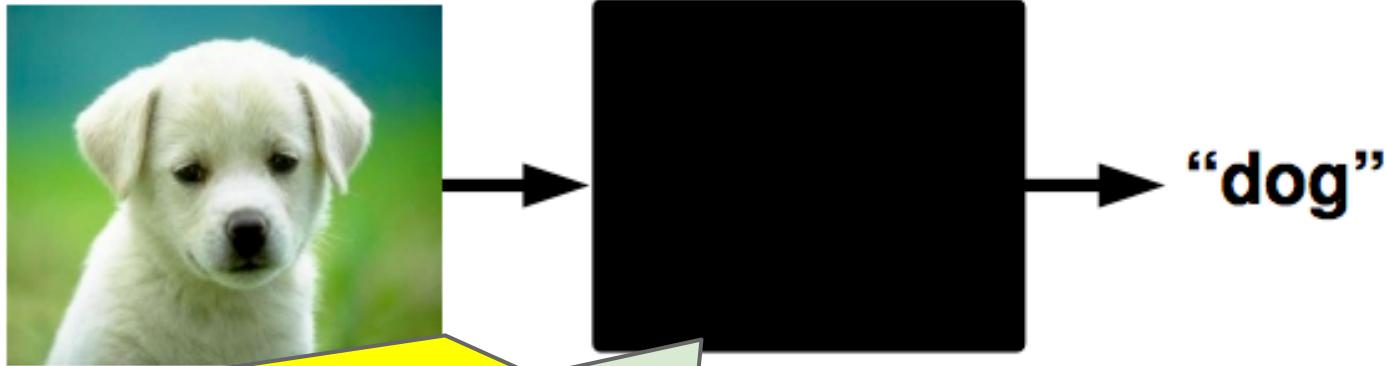
Preprocessing

Feature  
Extraction  
(HOG, SIFT, etc)

Post-processing  
(Feature selection,  
MKL etc)

Classifier  
(SVM,  
boosting, etc)

# The Black Box in a Traditional Recognition Approach



Preprocessing

Feature  
Extraction  
(HOG, SIFT, etc)

Post-processing  
(Feature selection,  
MKL etc)

- Most critical for accuracy
  - Most time-consuming in development
  - What is the best feature???
  - What is next?? Keep on crafting better features?
- ⇒ Let's learn feature representation directly from data.

# Learn features and classifier *together*

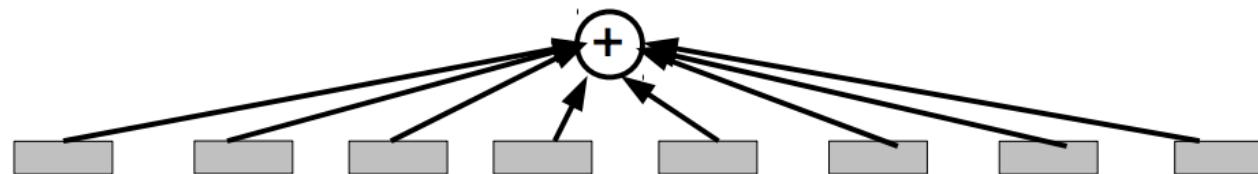
⇒ Learn an end-to-end recognition system.  
A non-linear map that takes raw pixels directly  
to labels.

**Q:** How can we build such a highly non-linear system?

**A:** By combining simple building blocks we can make more and more complex systems.

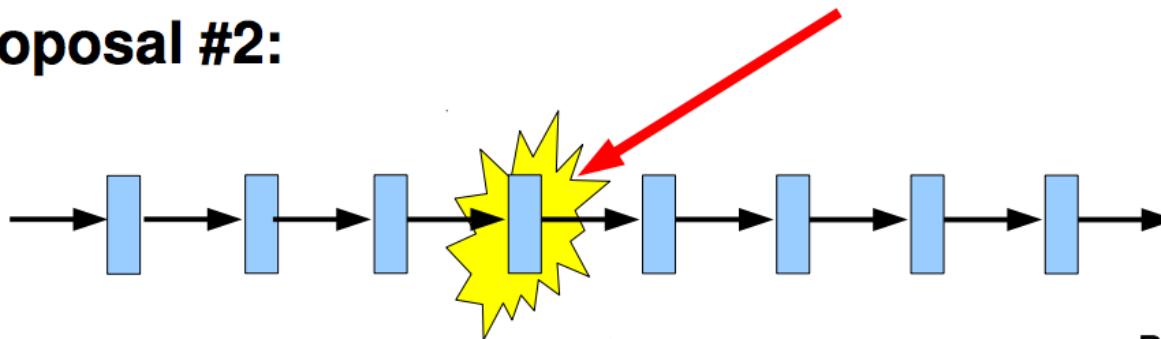
# Building a complicated function

**Proposal #1:**



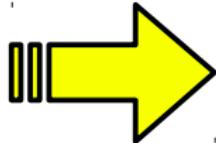
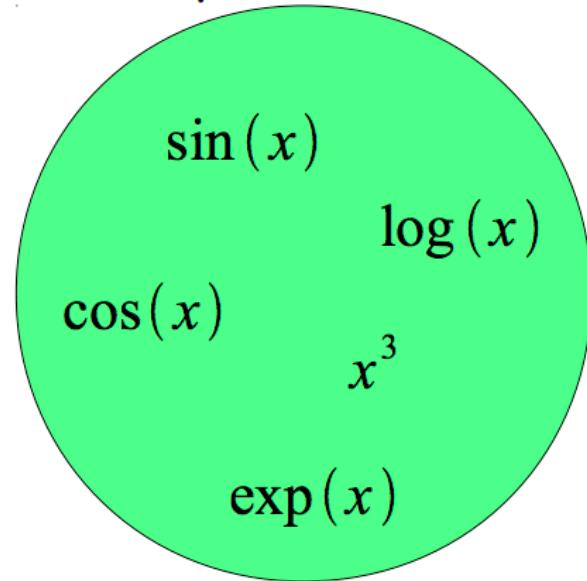
**Each box is a simple nonlinear function**

**Proposal #2:**



# Building a complicated function

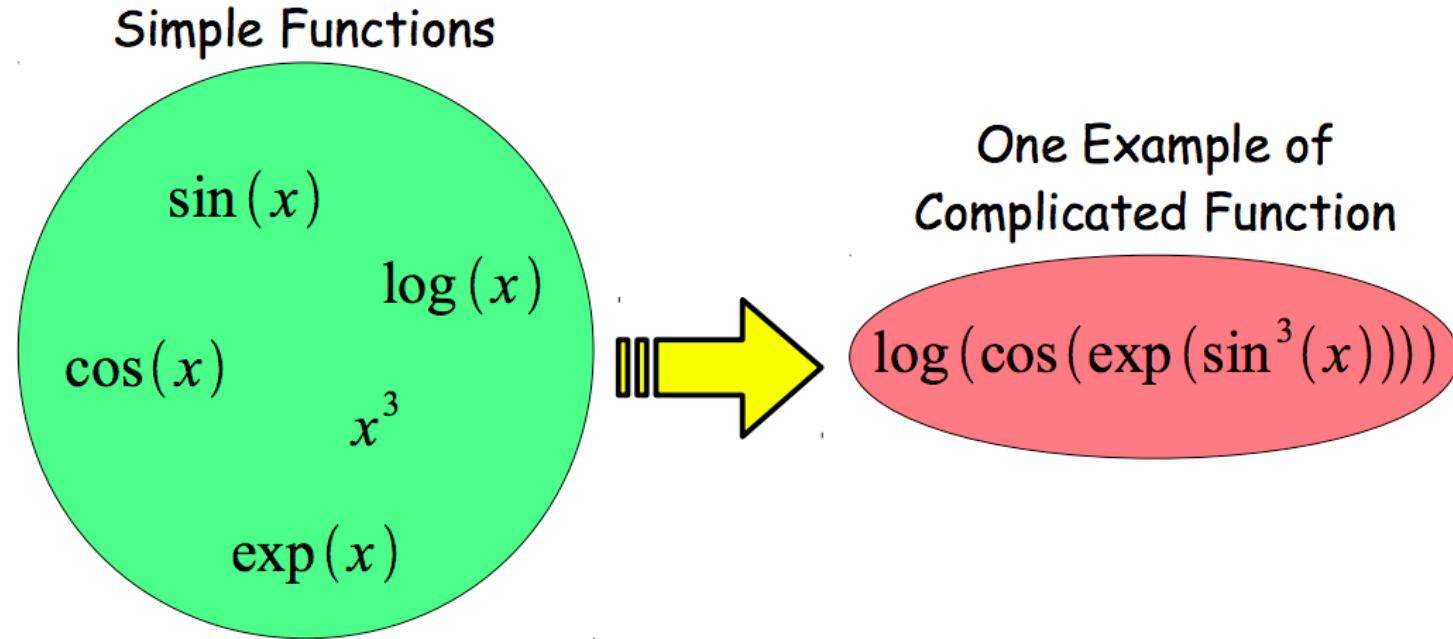
Simple Functions



One Example of  
Complicated Function

$\log(\cos(\exp(\sin^3(x))))$

# Building a complicated function

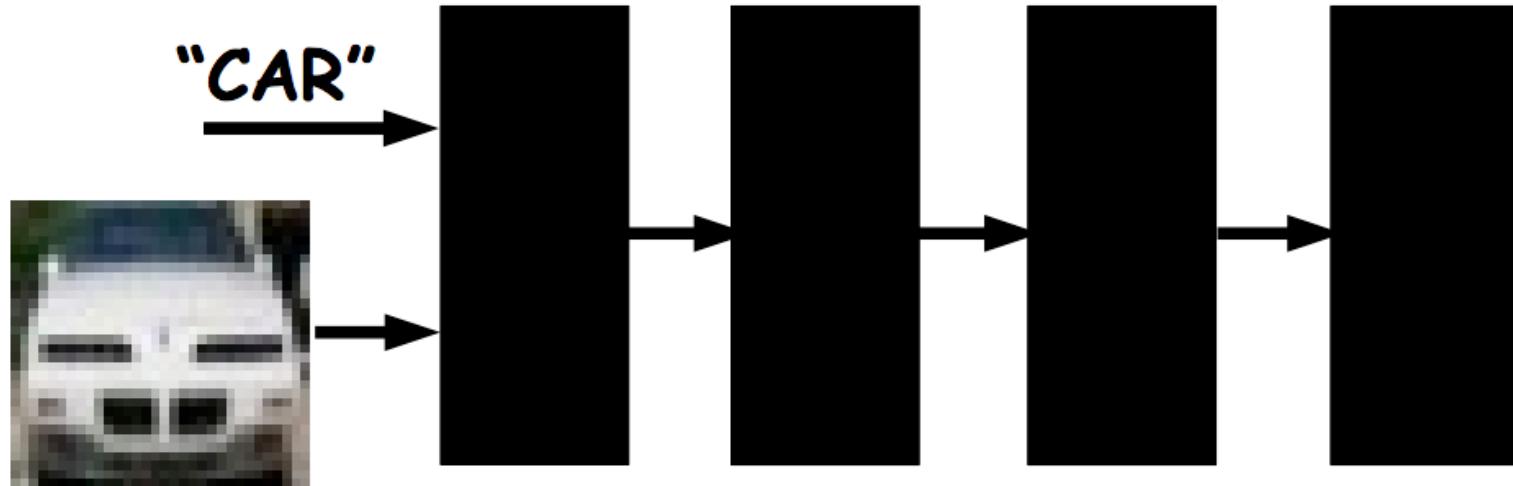


- Composition is at the core of deep learning methods
- Each “simple function” will have parameters subject to learning

# Intuition behind Deep Neural Nets

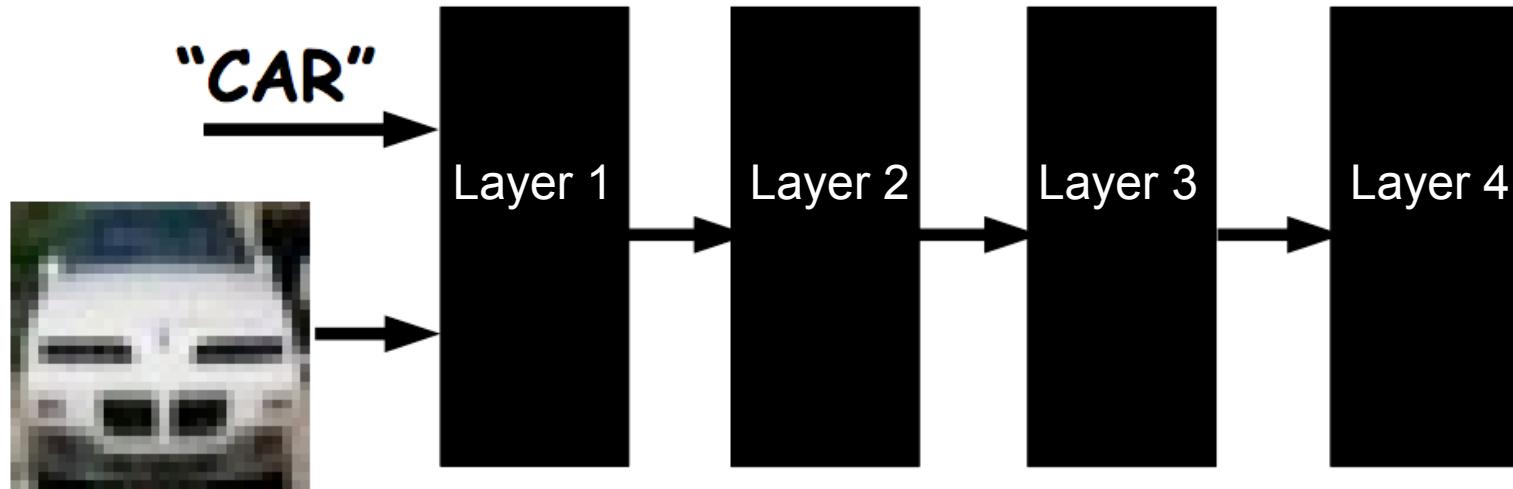


# Intuition behind Deep Neural Nets



**NOTE:** Each black box can have trainable parameters.  
Their composition makes a highly non-linear system.

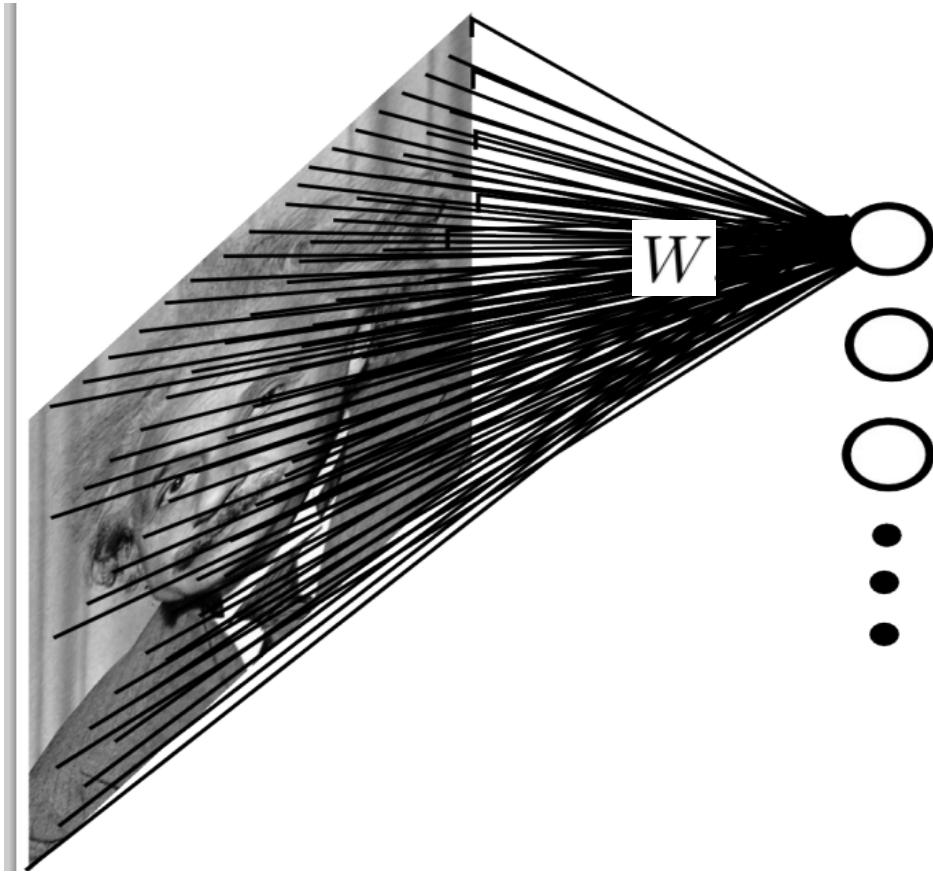
# Intuition behind Deep Neural Nets



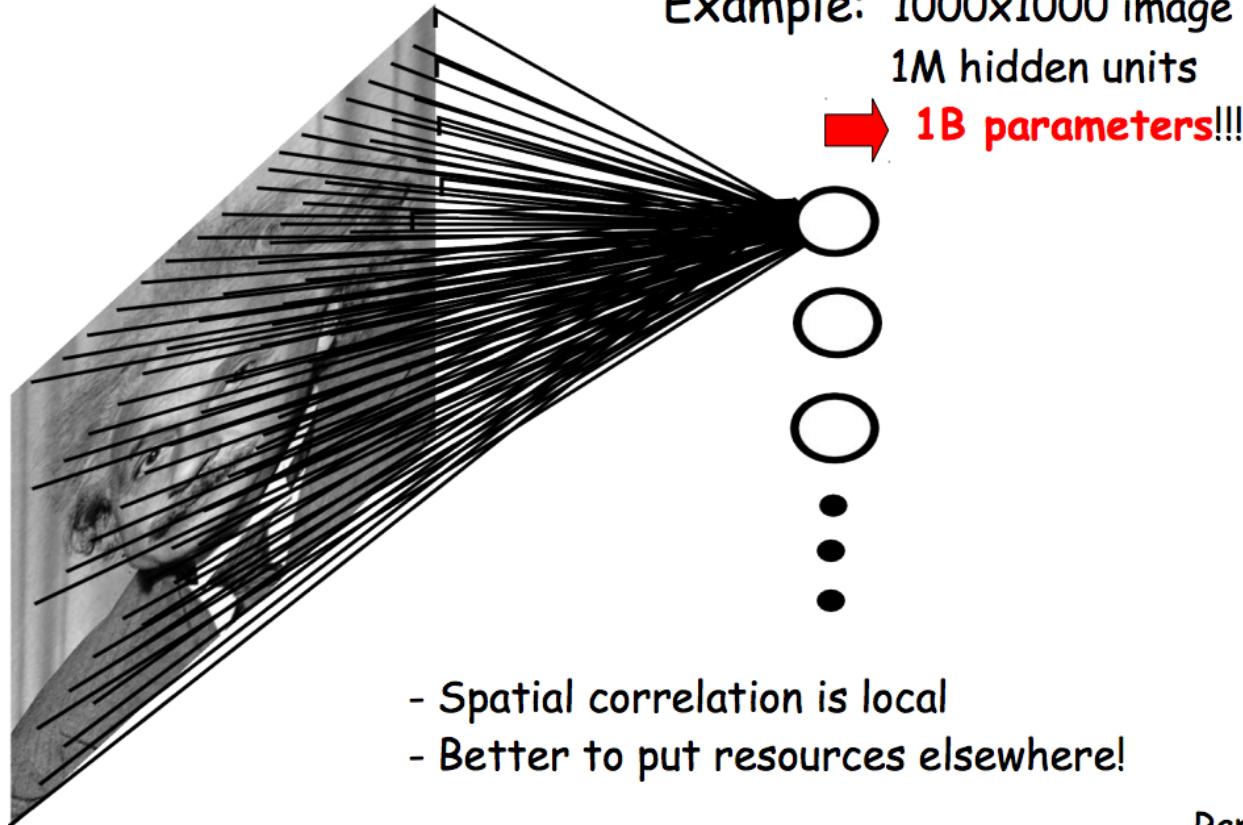
**NOTE:** Each black box can have trainable parameters.  
Their composition makes a highly non-linear system.

The final layer outputs a probability distribution of categories.

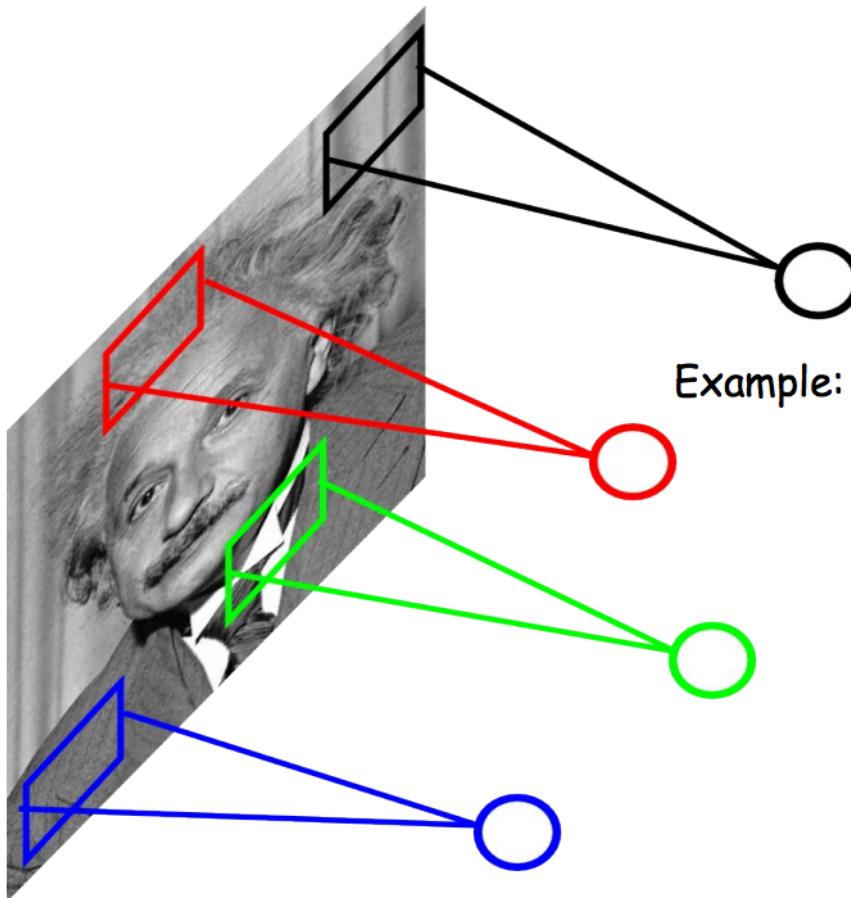
# When the input data is an image..



# When the input data is an image..

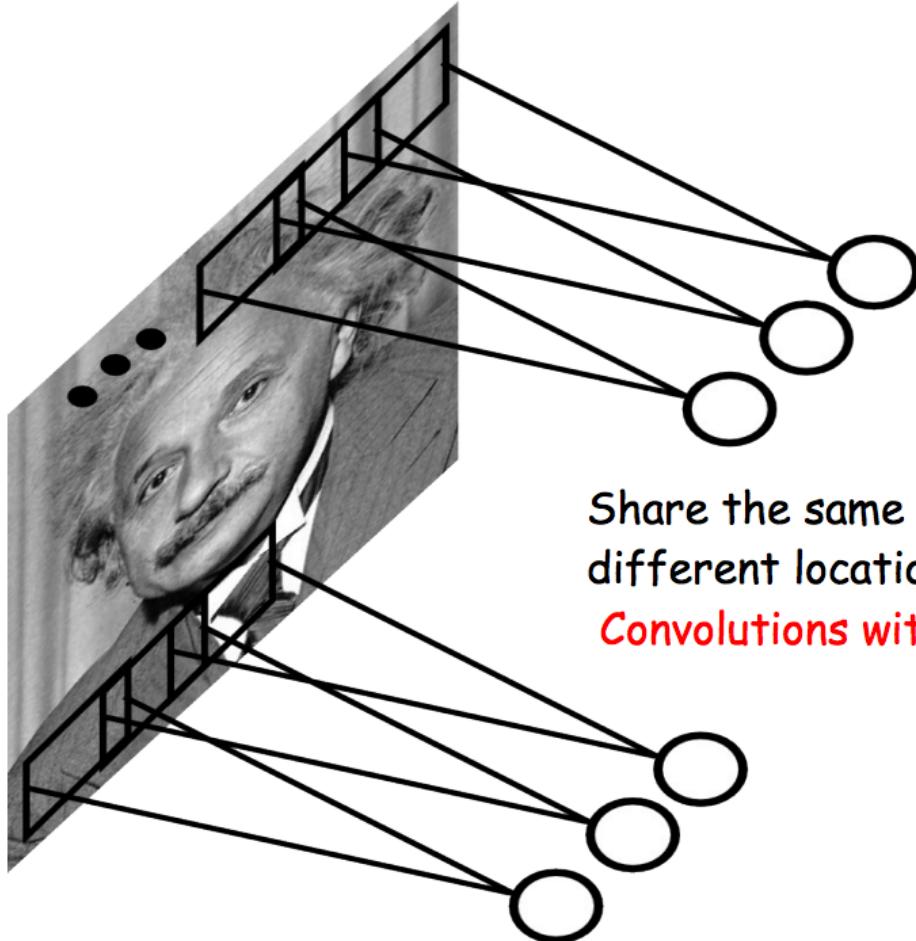


# Reduce connection to local regions



Example: 1000x1000 image  
1M hidden units  
Filter size: 10x10  
10M parameters

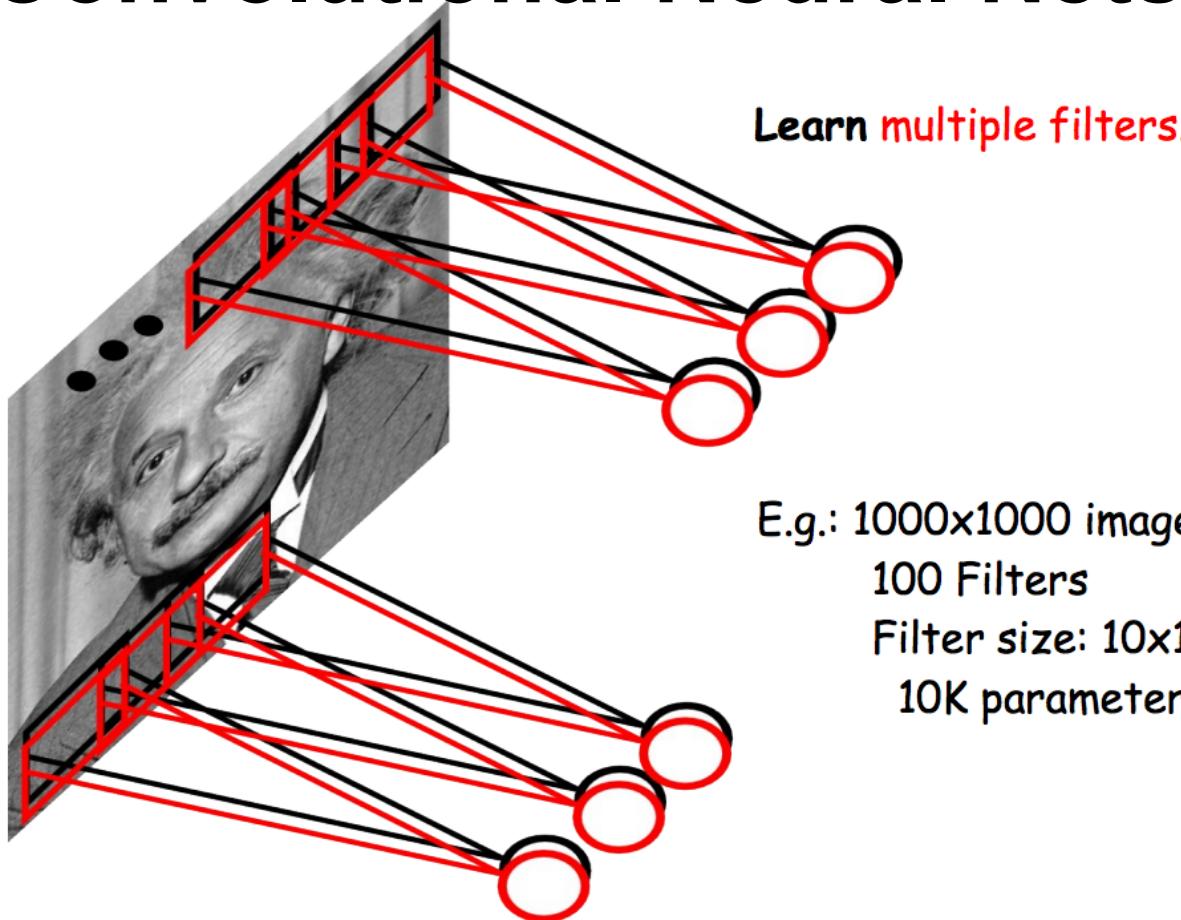
# Reuse the same kernel everywhere



Because interesting features (edges) can happen at anywhere in the image.

Share the same parameters across  
different locations:  
**Convolutions with learned kernels**

# Convolutional Neural Nets



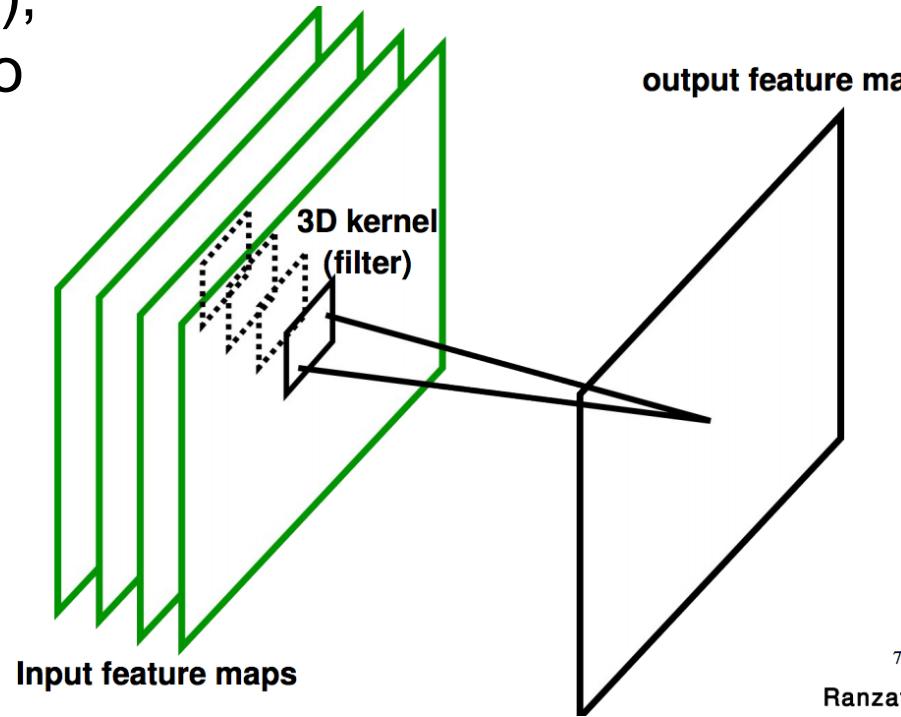
LeCun et al. "Gradient-based learning applied to document recognition" IEEE 1998

# Detail

If the input has 3 channels (R,G,B),  
3 separate  $k$  by  $k$  filter is applied to  
each channel.

Output of convolving 1 feature is  
called a *feature map*.

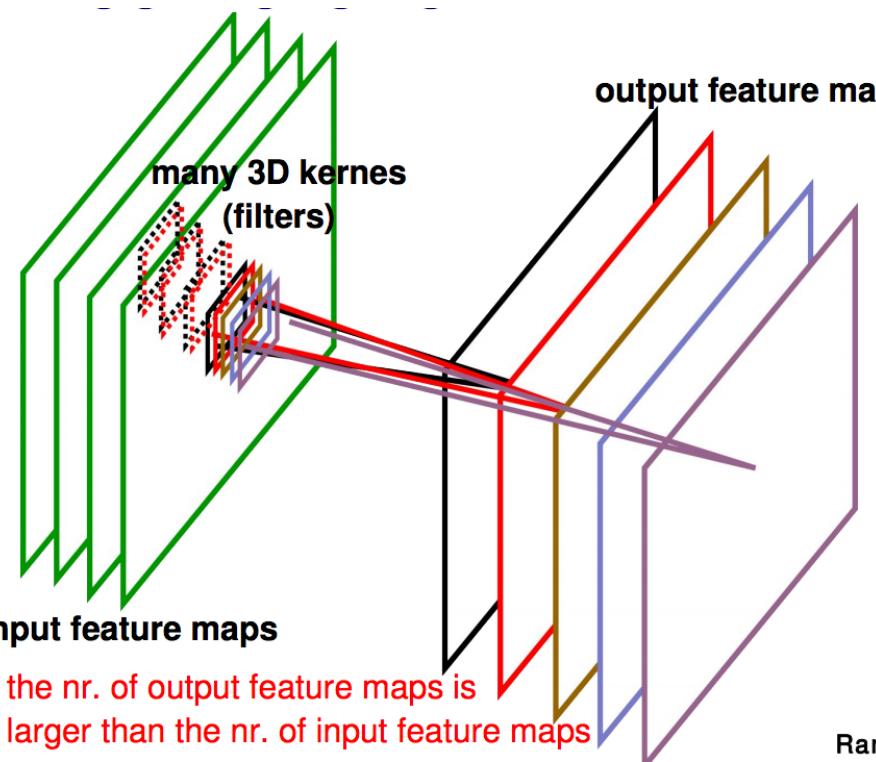
This is just sliding window, ex. the output of  
one part filter of DPM is a feature map



# Using multiple filters

Each filter detects features in the output of previous layer.

So to capture different features, learn multiple filters.



NOTE: the nr. of output feature maps is  
usually larger than the nr. of input feature maps

# Example of filtering

- Convolutional
  - Translation equivariance
  - Tied filter weights  
(same at each position → few parameters)

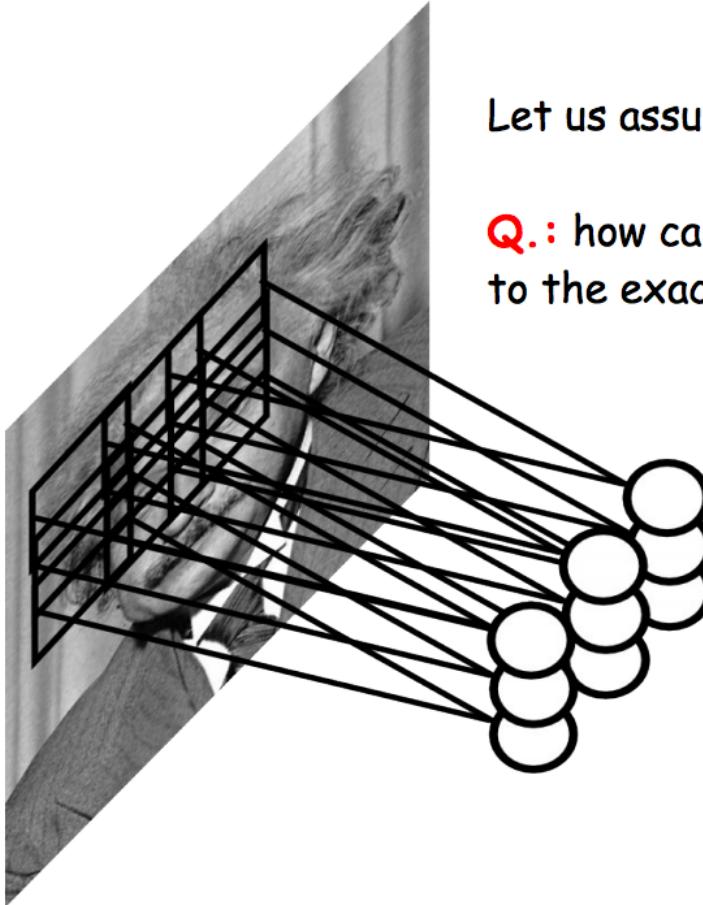


Input



Feature Map

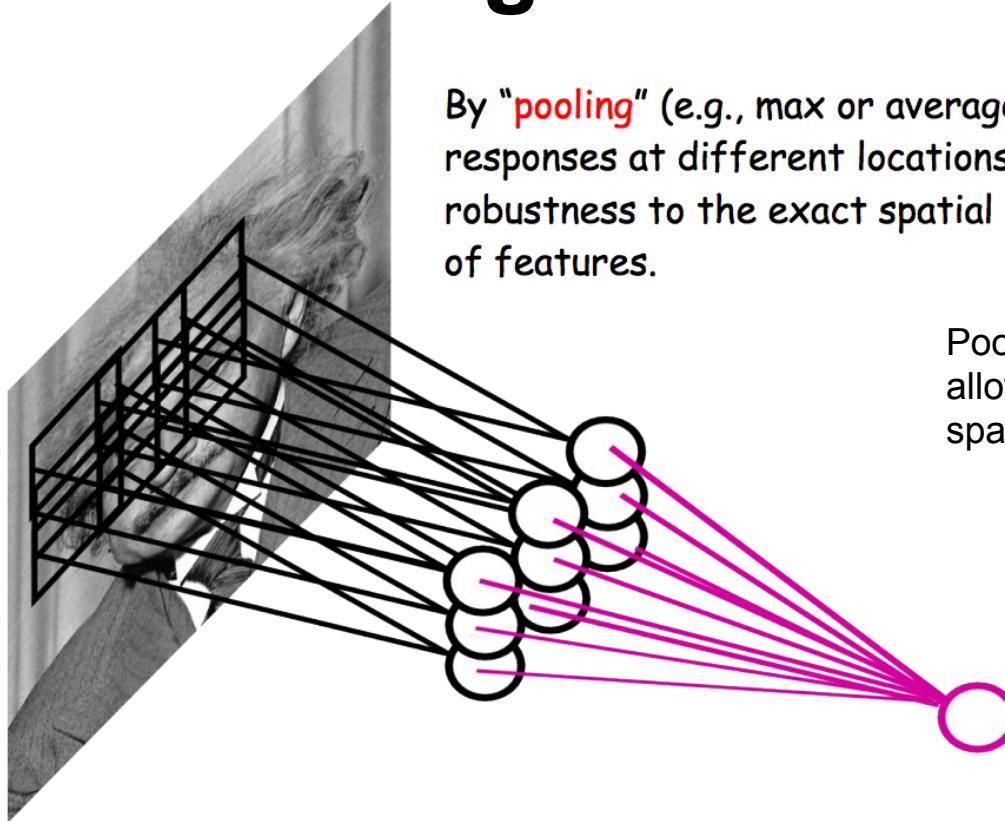
# Building Translation Invariance



Let us assume filter is an "eye" detector.

**Q.:** how can we make the detection robust  
to the exact location of the eye?

# Building Translation Invariance via Spatial Pooling



By “**pooling**” (e.g., max or average) filter responses at different locations we gain robustness to the exact spatial location of features.

Pooling also subsamples the image, allowing the next layer to look at larger spatial regions.

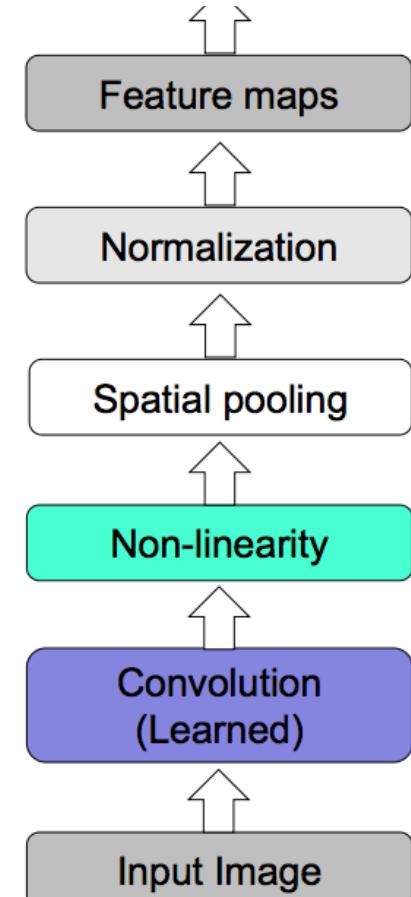
# Summary of a typical convolutional layer

Doing all of this consists one layer.

- Pooling and normalization is optional.

Stack them up and train just like multi-layer neural nets.

Final layer is usually fully connected neural net with output size == number of classes



# Revisiting the composition idea

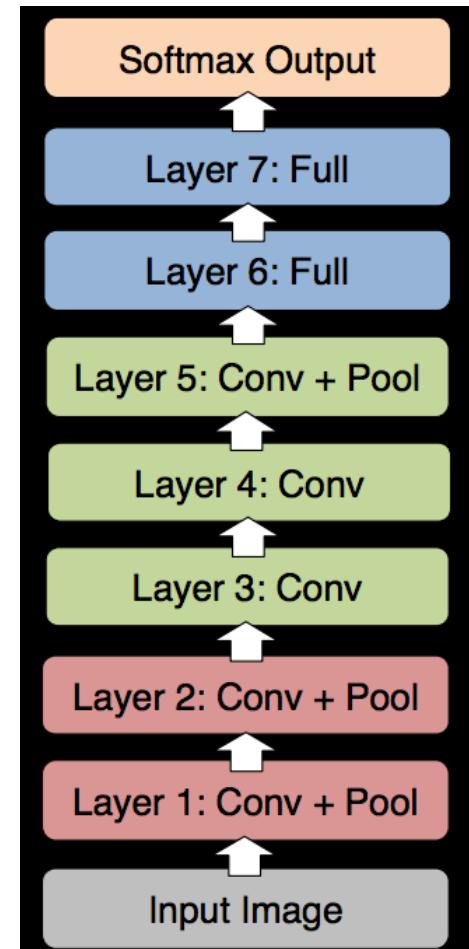
Every layer learns a feature detector by combining the output of the layer before.

⇒ More and more abstract features are learned as we stack layers.

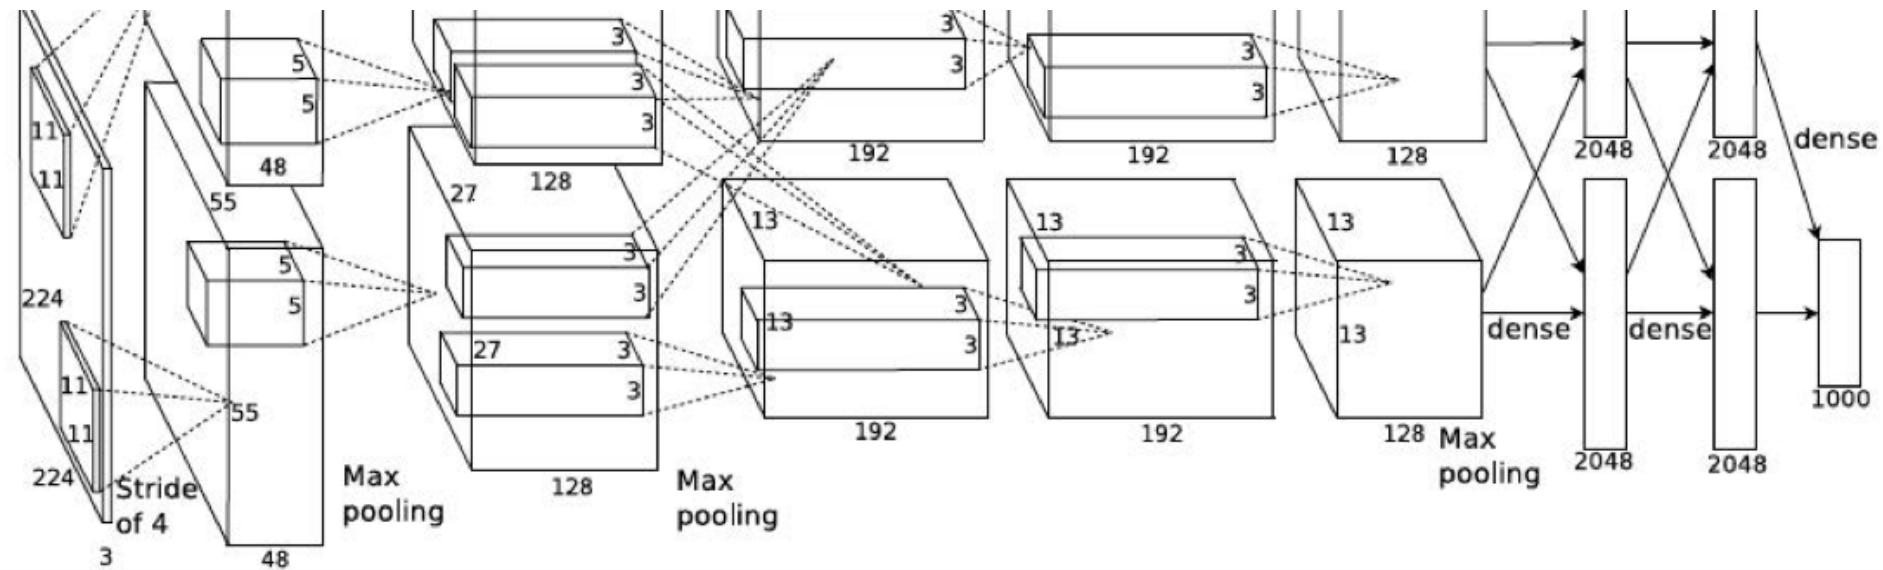
Keep this in mind and let's look at what kind of things ConvNets learn.

# Architecture of Alex Krizhevsky et al.

- 8 layers total.
- Trained on Imagenet Dataset (1000 categories, 1.2M training images, 150k test images)
- 18.2% top-5 error
  - Winner of the ILSVRC-2012 challenge.



# Architecture of Alex Krizhevsky et al.

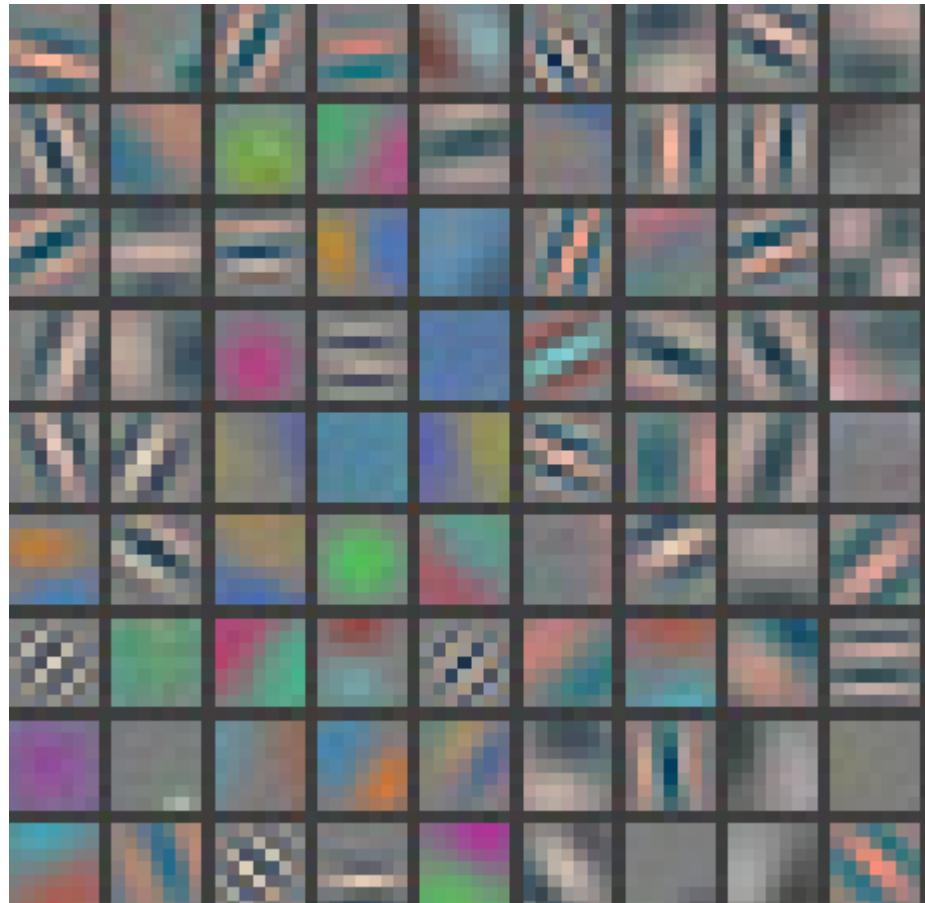


# First layer filters

Showing 81 filters of  
11x11x3.

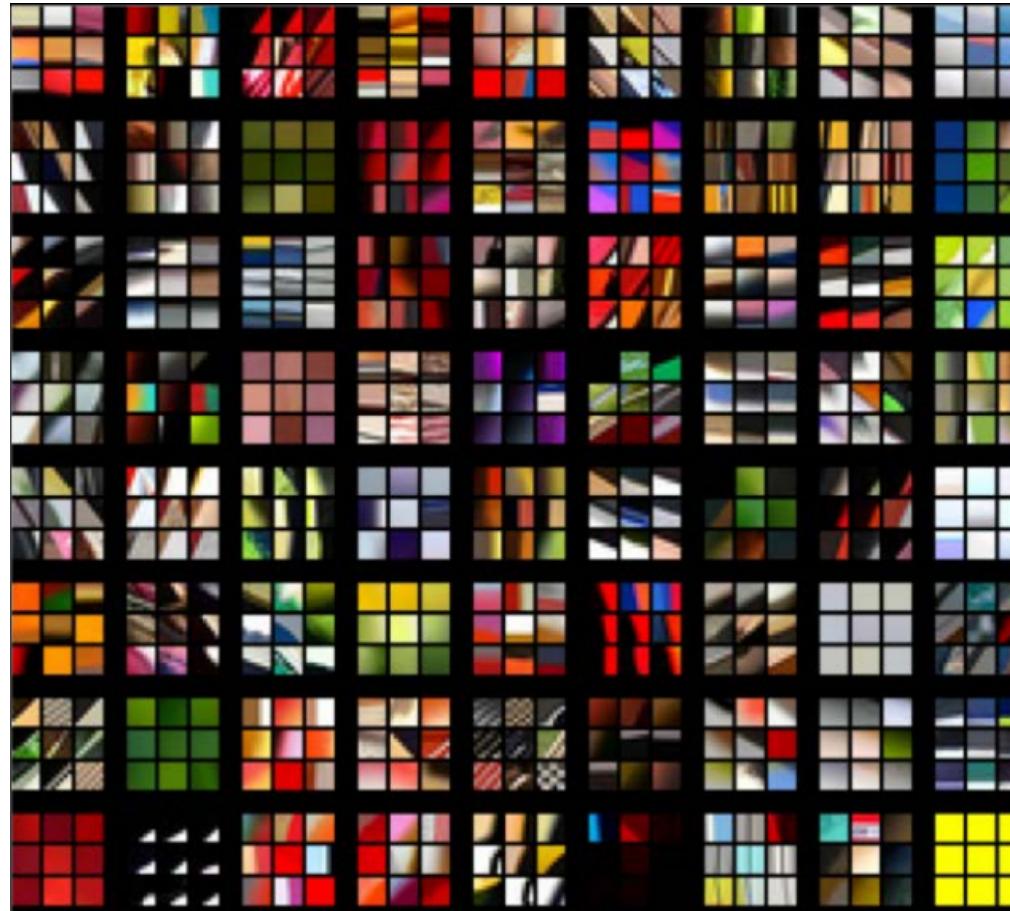
Capture low-level  
features like oriented  
edges, blobs.

Note these oriented edges are  
analogous to what SIFT uses to  
compute the gradients.

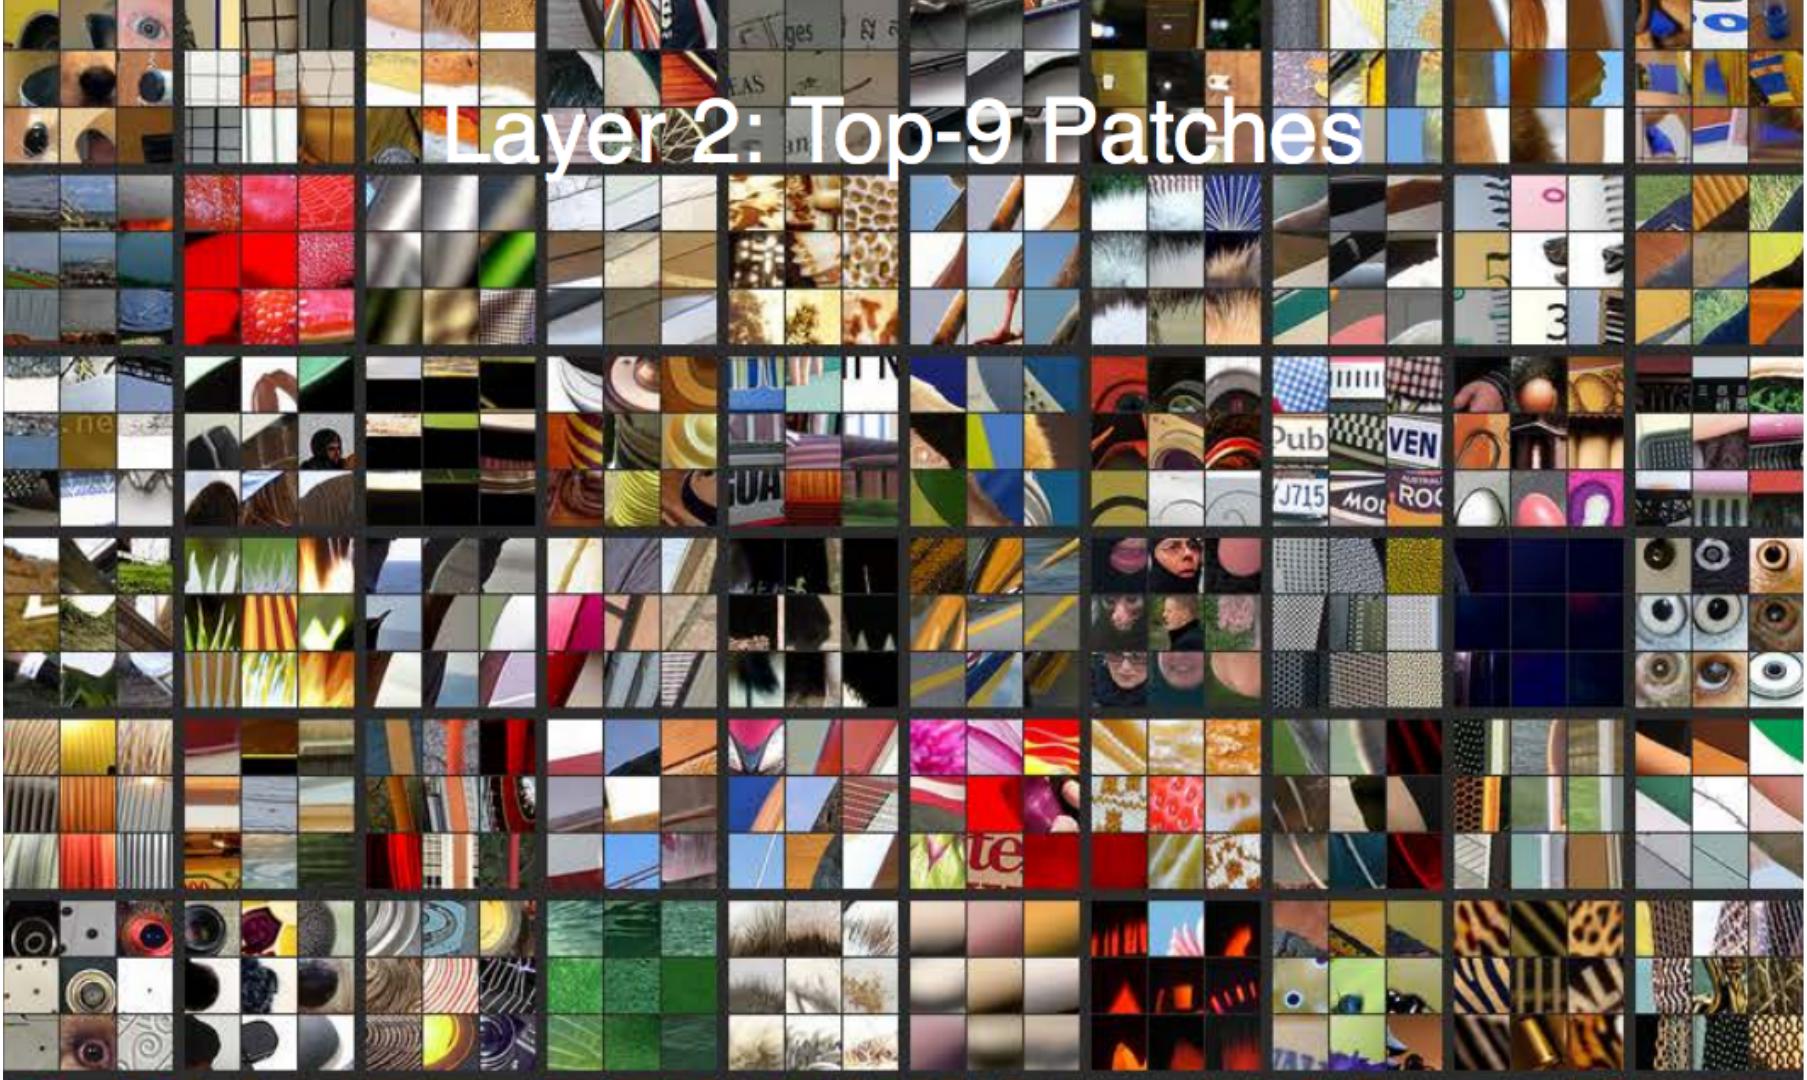


# Top 9 patches that activate each filter in layer 1

Each 3x3 block shows  
the top 9 patches for  
one filter.



# Layer 2: Top-9 Patches





## Layer 2: Top-9 Patches

Note how the previous low-level features are combined to detect a little more abstract features like textures.

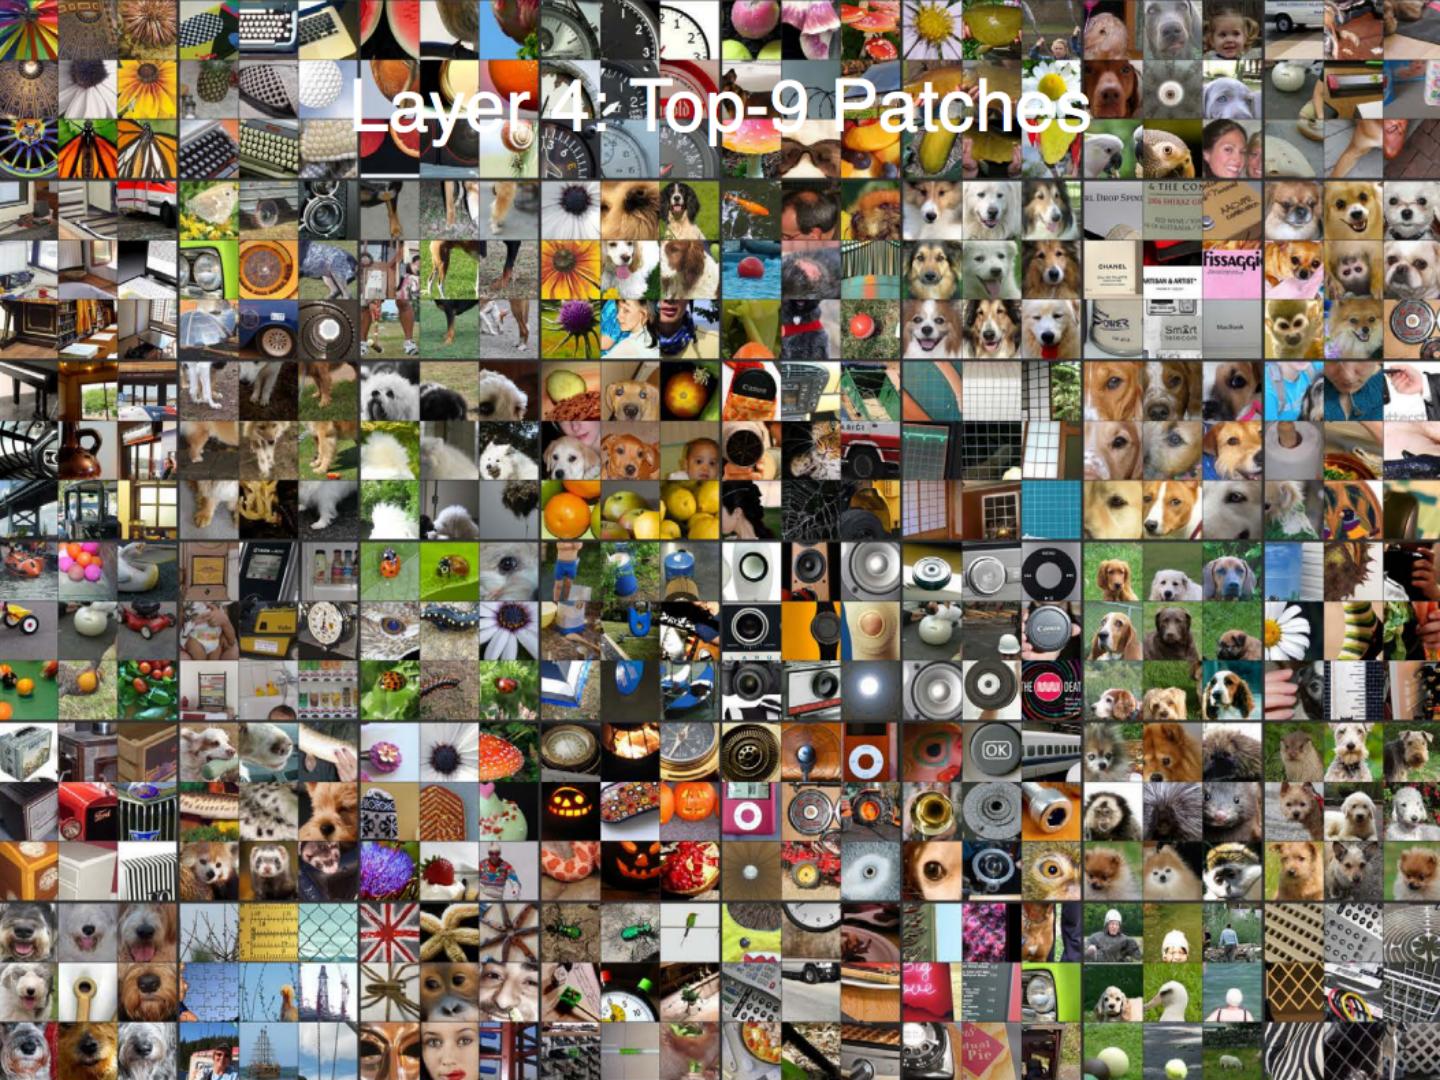
# Layer 3: Top-9 Patches



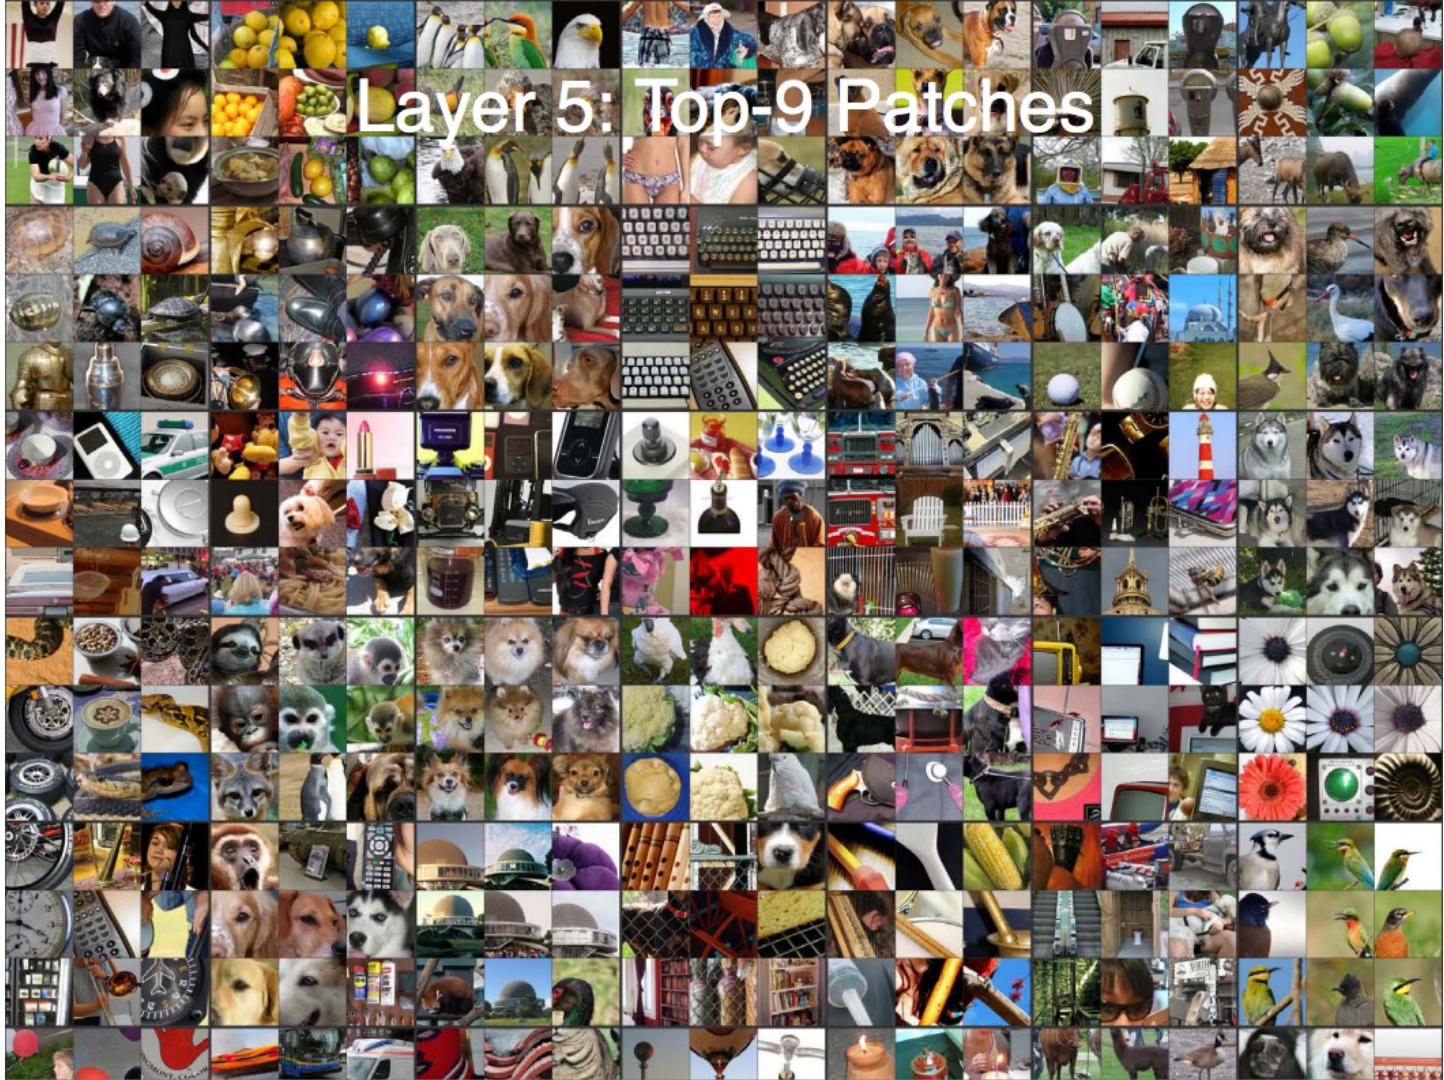
# Layer 3: Top-9 Patches



# Layer 4: Top-9 Patches



# Layer 5: Top-9 Patches



# ConvNets as generic feature extractor

- A well-trained ConvNets is an excellent feature extractor.
- Chop the network at desired layer and use the output as a feature representation to train a SVM on some other vision dataset.

	Cal-101 (30/class)	Cal-256 (60/class)
SVM (1)	$44.8 \pm 0.7$	$24.6 \pm 0.4$
SVM (2)	$66.2 \pm 0.5$	$39.6 \pm 0.3$
SVM (3)	$72.3 \pm 0.4$	$46.0 \pm 0.3$
SVM (4)	$76.6 \pm 0.4$	$51.3 \pm 0.1$
SVM (5)	<b><math>86.2 \pm 0.8</math></b>	$65.6 \pm 0.3$
SVM (7)	<b><math>85.5 \pm 0.4</math></b>	<b><math>71.7 \pm 0.2</math></b>
Softmax (5)	$82.9 \pm 0.4$	$65.7 \pm 0.5$
Softmax (7)	<b><math>85.4 \pm 0.4</math></b>	<b><math>72.6 \pm 0.1</math></b>

- Improve further by taking a pre-trained ConvNet and re-training it on a different dataset. Called *fine-tuning*